



# **ExplosION: The Hidden Mines in the Android ION Driver**

Le Wu, Xuen Li, Tim Xia

Baidu Security

# About us

## Le Wu([@Nvamous](#))

- Focus on Android/Linux bug hunting and exploit
- Found 200+ vulnerabilities in the last two years
- Top1 in Android Chipset Security Program, Top1 in MediaTek Mobile Security Program

## Xuen Li([@lxn524](#))

- Interested in Android, Linux Kernel and OSS security testing and exploitation
- Focus on PSA and OP-TEE currently
- Author of One Click Root Master, SuperRootMaster, etc...

## Tim Xia

- Staff security researcher at Baidu Inc.
- Focused on system and software security solutions
- Previous PacSec and HITB speaker

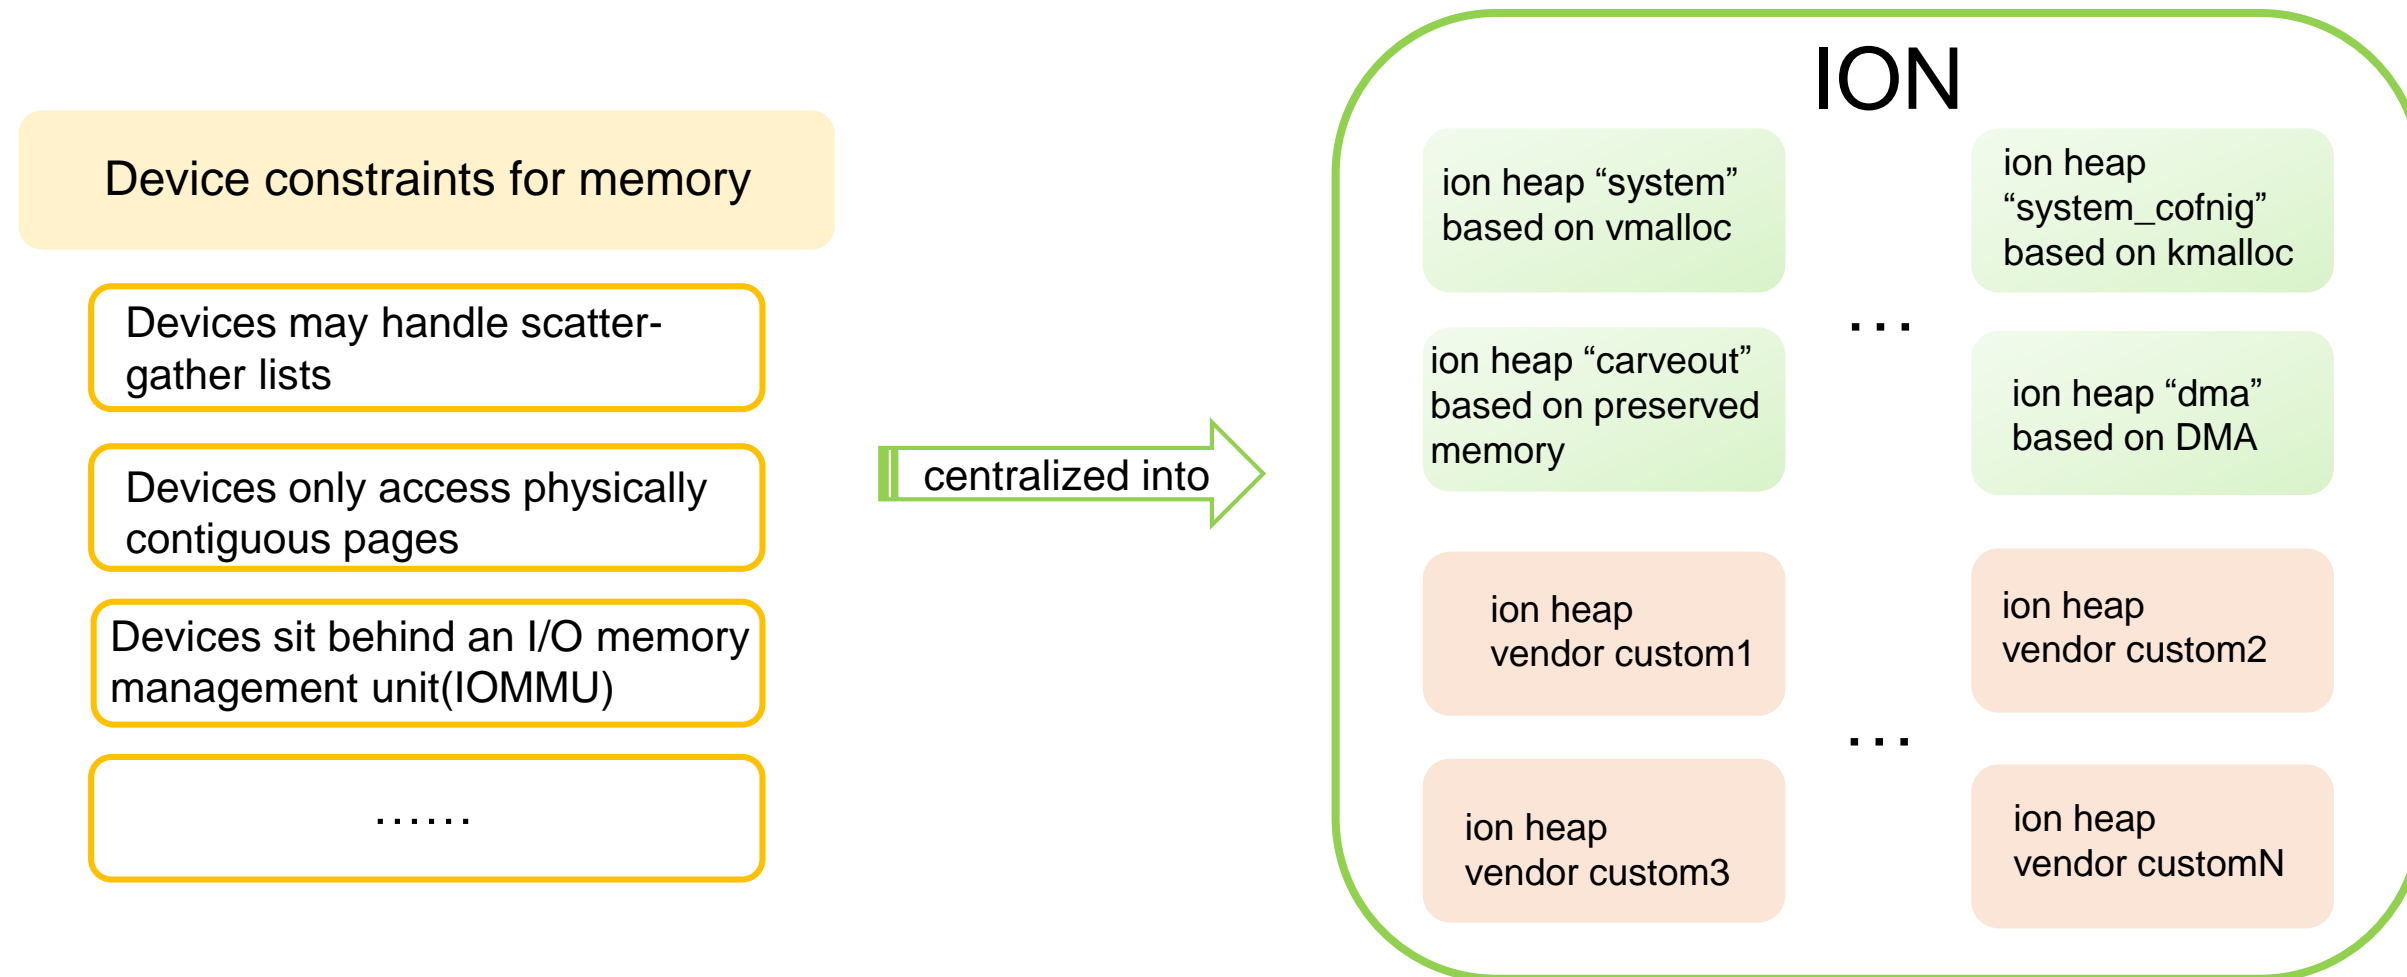
# Agenda

- Introduction to ION
- Using ION
- Diving into ExplosION
- Reflections on ExplosION
- Future work

# Introduction—What is ION?

- A generalized memory pool manager

It's used to address the issue of fragmented memory management interfaces across different Android devices.



# Introduction——Why ION?

- A common Android driver used for a decade
  - ✓ Although ION is replaced with dma-buf in Android 12, millions of devices still use ION
  - ✓ ION will survive for a long time due to fragmentation of Android
- Can be accessed by untrusted apps
  - ✓ No additional privileges needed to trigger the vulnerabilities in ION
  - ✓ Perfect exploitation target, like the BINDER device

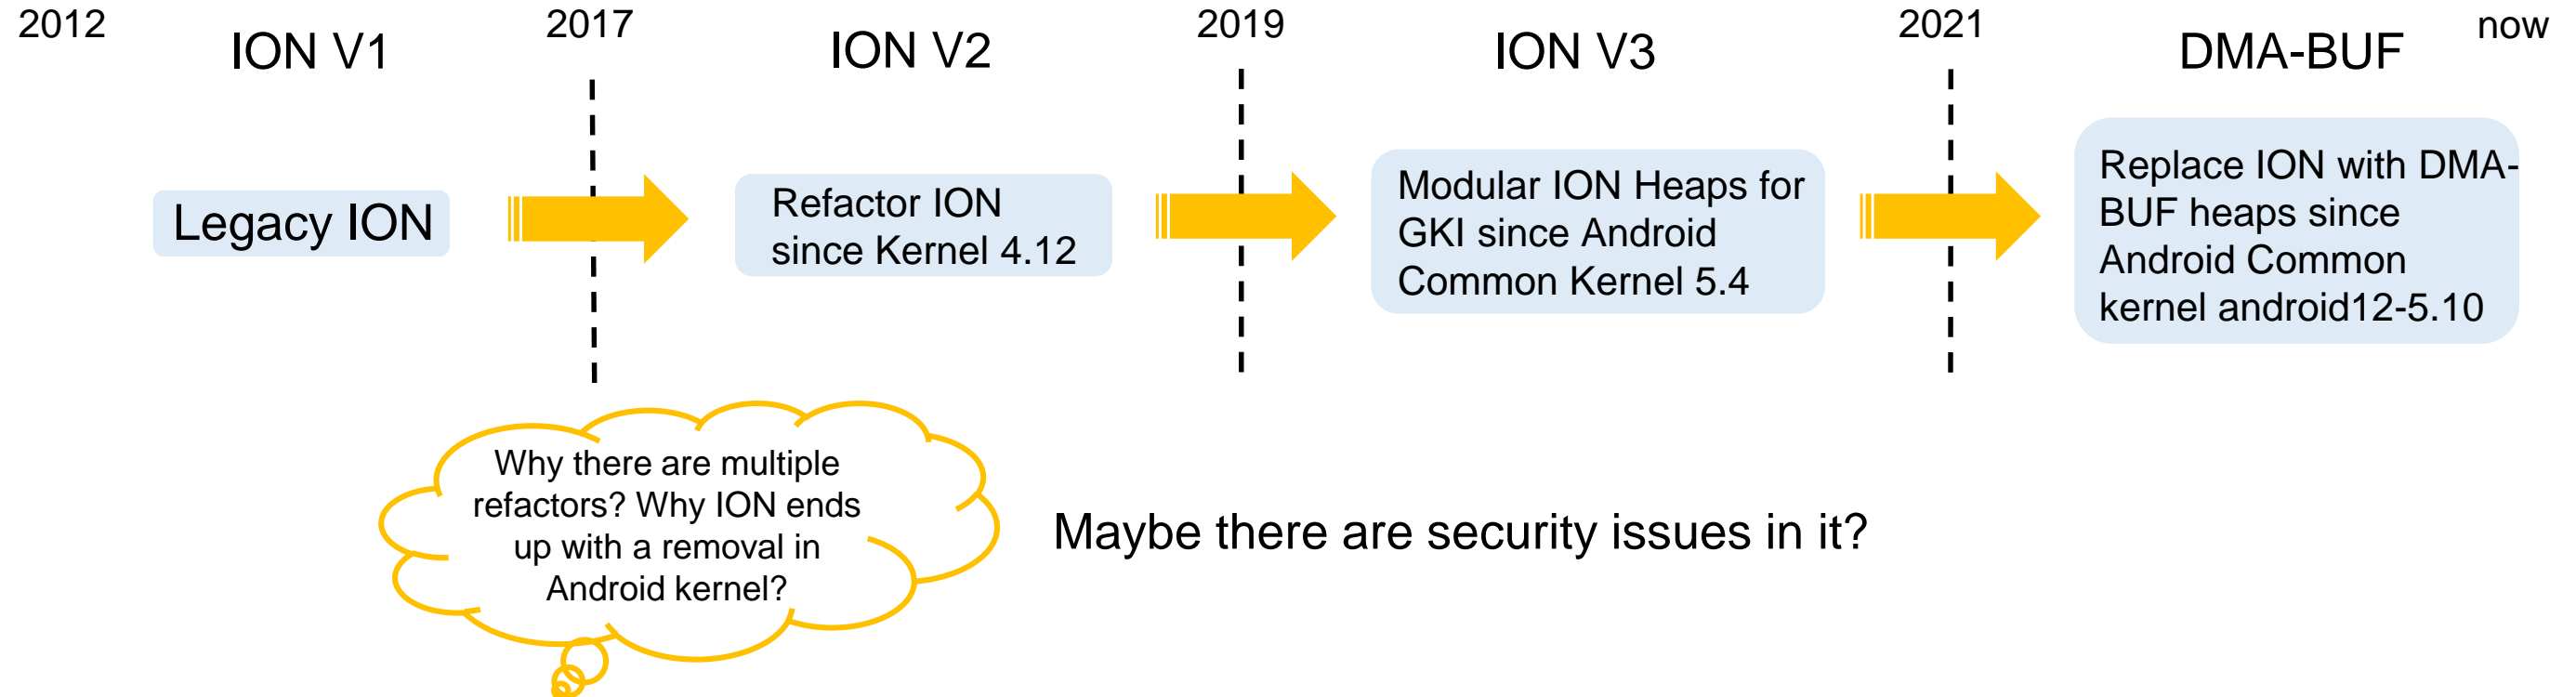
# Introduction——Why ION?

- A memory management component

Vulnerabilities in memory management components are proven to be dangerous!



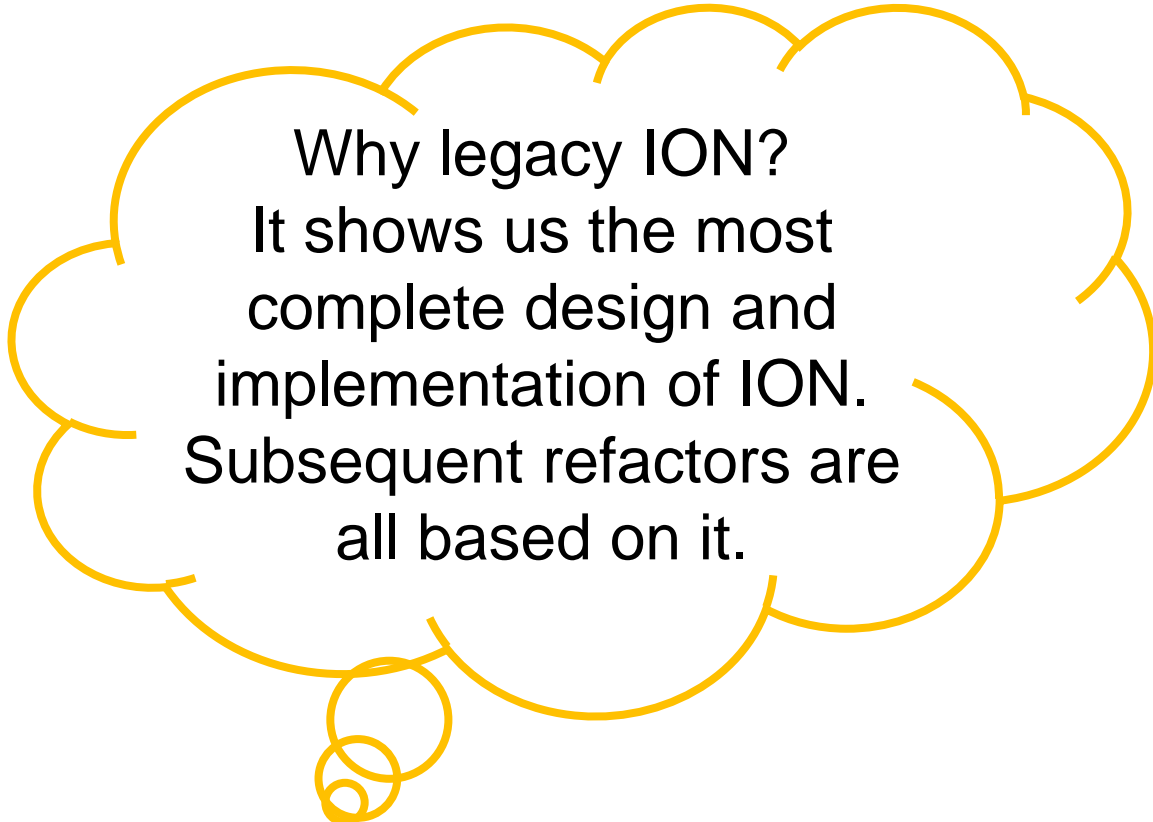
- Multiple refactors from upstream



So far, we have found 40+ vulnerabilities in ION, and millions of devices are affected!



Before diving into ExplosION, let's have a look at the using of legacy ION



Why legacy ION?  
It shows us the most  
complete design and  
implementation of ION.  
Subsequent refactors are  
all based on it.



# Using ION

## ● Allocate memory buffer from user space

User space

```
client_fd = open(/dev/ion, RONLY);
```

```
ioctl(client_fd, ION_IOC_ALLOC, struct  
allocation_data *allocation_data)
```

ion\_handle id

```
Struct ion_allocation_data {  
    size_t len;  
    size_t align;  
    unsigned int heap_id_mask;  
    unsigned int flags;  
    int handle; // ion_handle id  
}
```

Kernel space

ion\_client

struct rb\_root handles

ion\_handle

ion\_handle

ion\_handle

.....

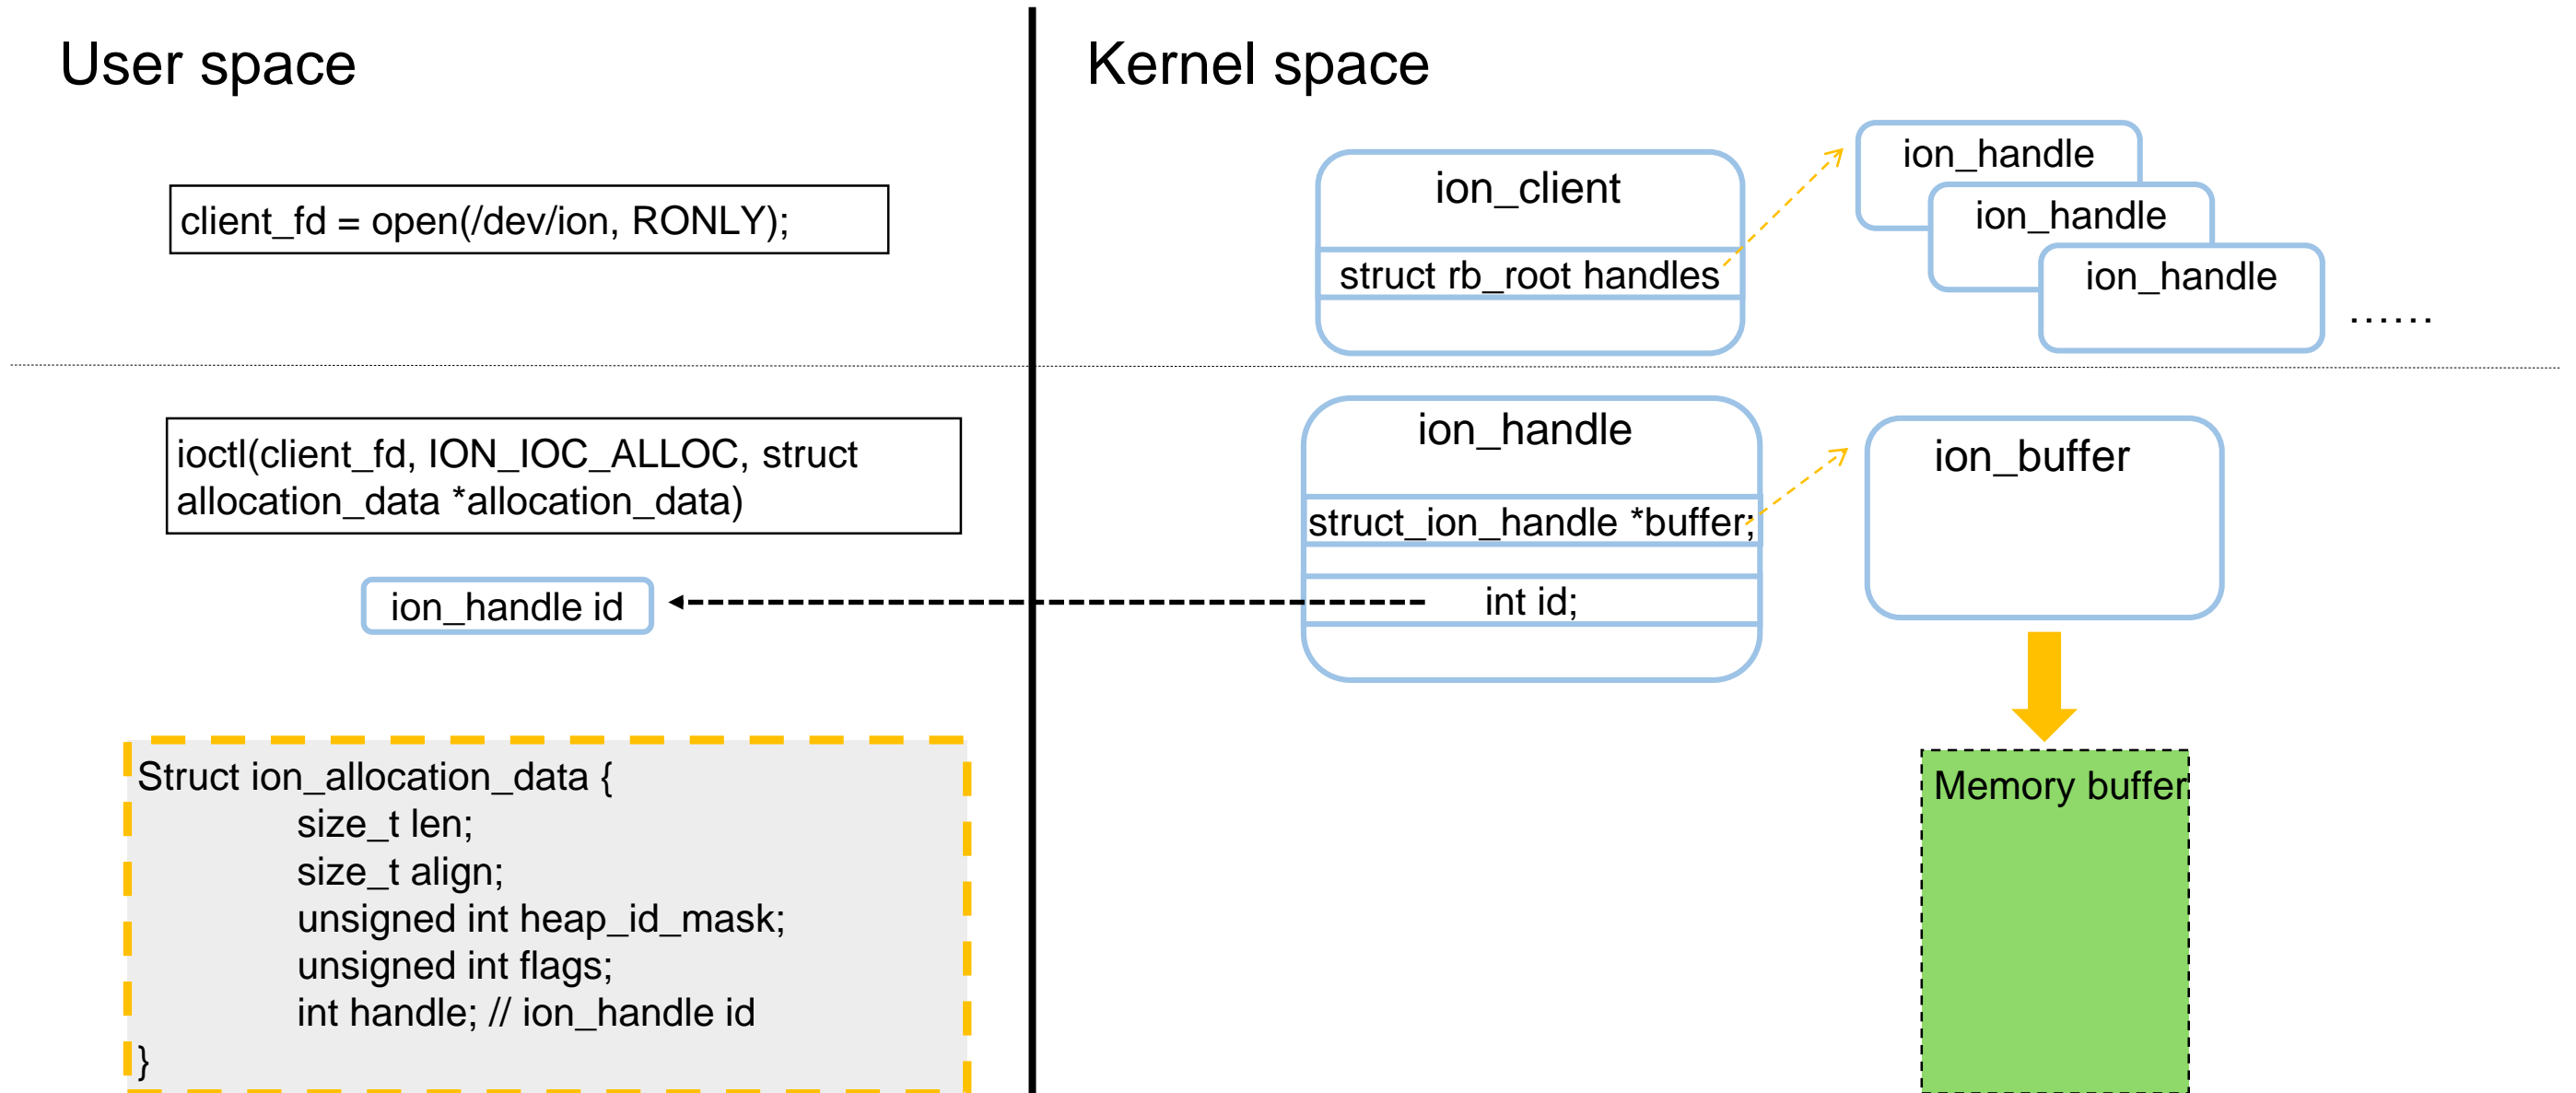
ion\_handle

struct\_ion\_handle \*buffer;

int id;

ion\_buffer

Memory buffer



# Using ION

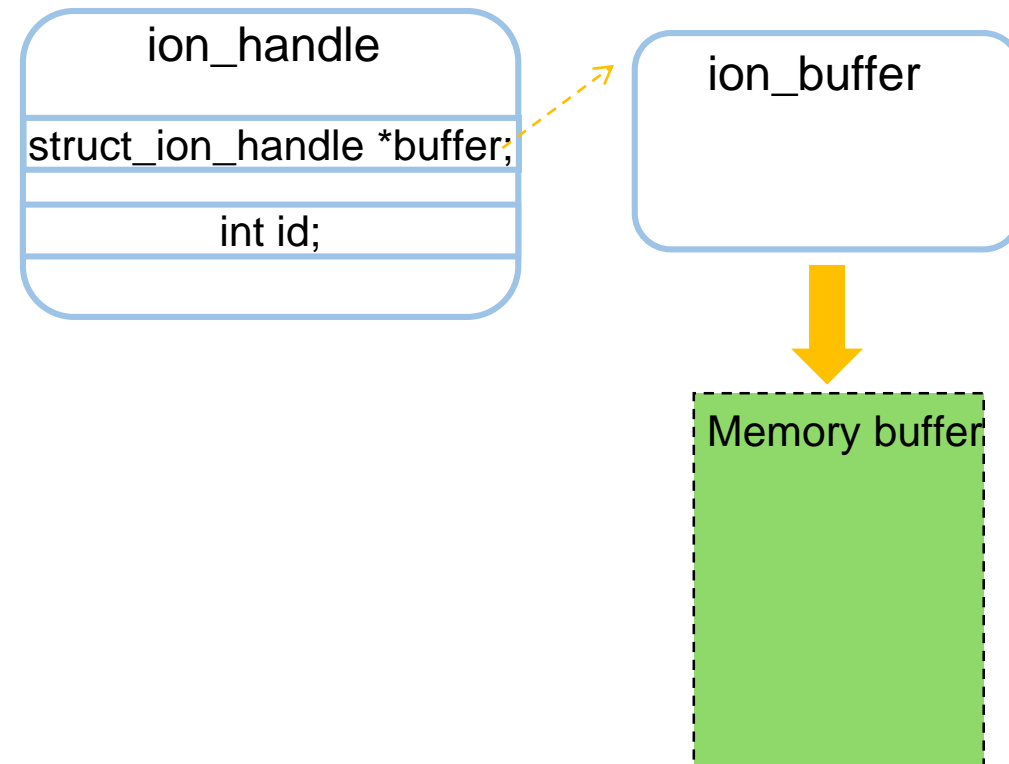
- Free memory buffer from user space

User space

```
ioctl(client_fd, ION_IOC_FREE, struct  
ion_handle_data *handle_data)
```

```
Struct ion_handle_data {  
    int handle; // ion_handle id  
}
```

Kernel space



# Using ION

## ● Buffer sharing from user space

User space

```
ioctl(client_fd, ION_IOC_SHARE, struct  
ion_fd_data *fd_data)
```

dma-buf fd

```
Struct ion_fd_data {  
    int handle; // input:ion_handle fd  
    int fd; // output:dma-buf fd  
}
```

```
mmap(dma-buf fd,.....)
```

Memory buffer

Kernel space

ion\_handle

struct\_ion\_handle \*buffer;

int id;

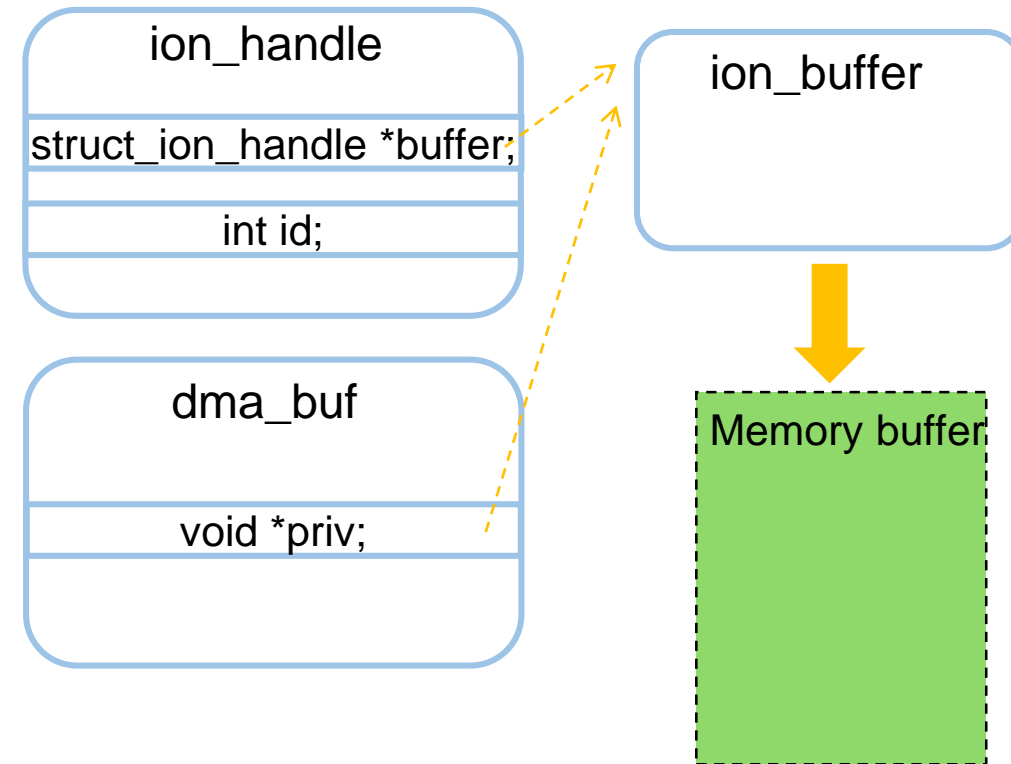
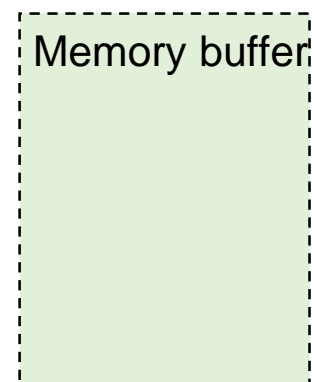
dma\_buf

void \*priv;

dma\_buf\_fd()

ion\_buffer

Memory buffer



# Using ION

## ● Import buffer from user space

User space

```
ioctl(client_fd, ION_IOC_IMPORT,  
struct ion_fd_data *fd_data)
```

ion\_handle id

```
Struct ion_fd_data {  
    int handle; //output:ion_handle fd  
    int fd; // input:dma-buf fd  
}
```

Kernel space

dma\_buf

void \*priv;

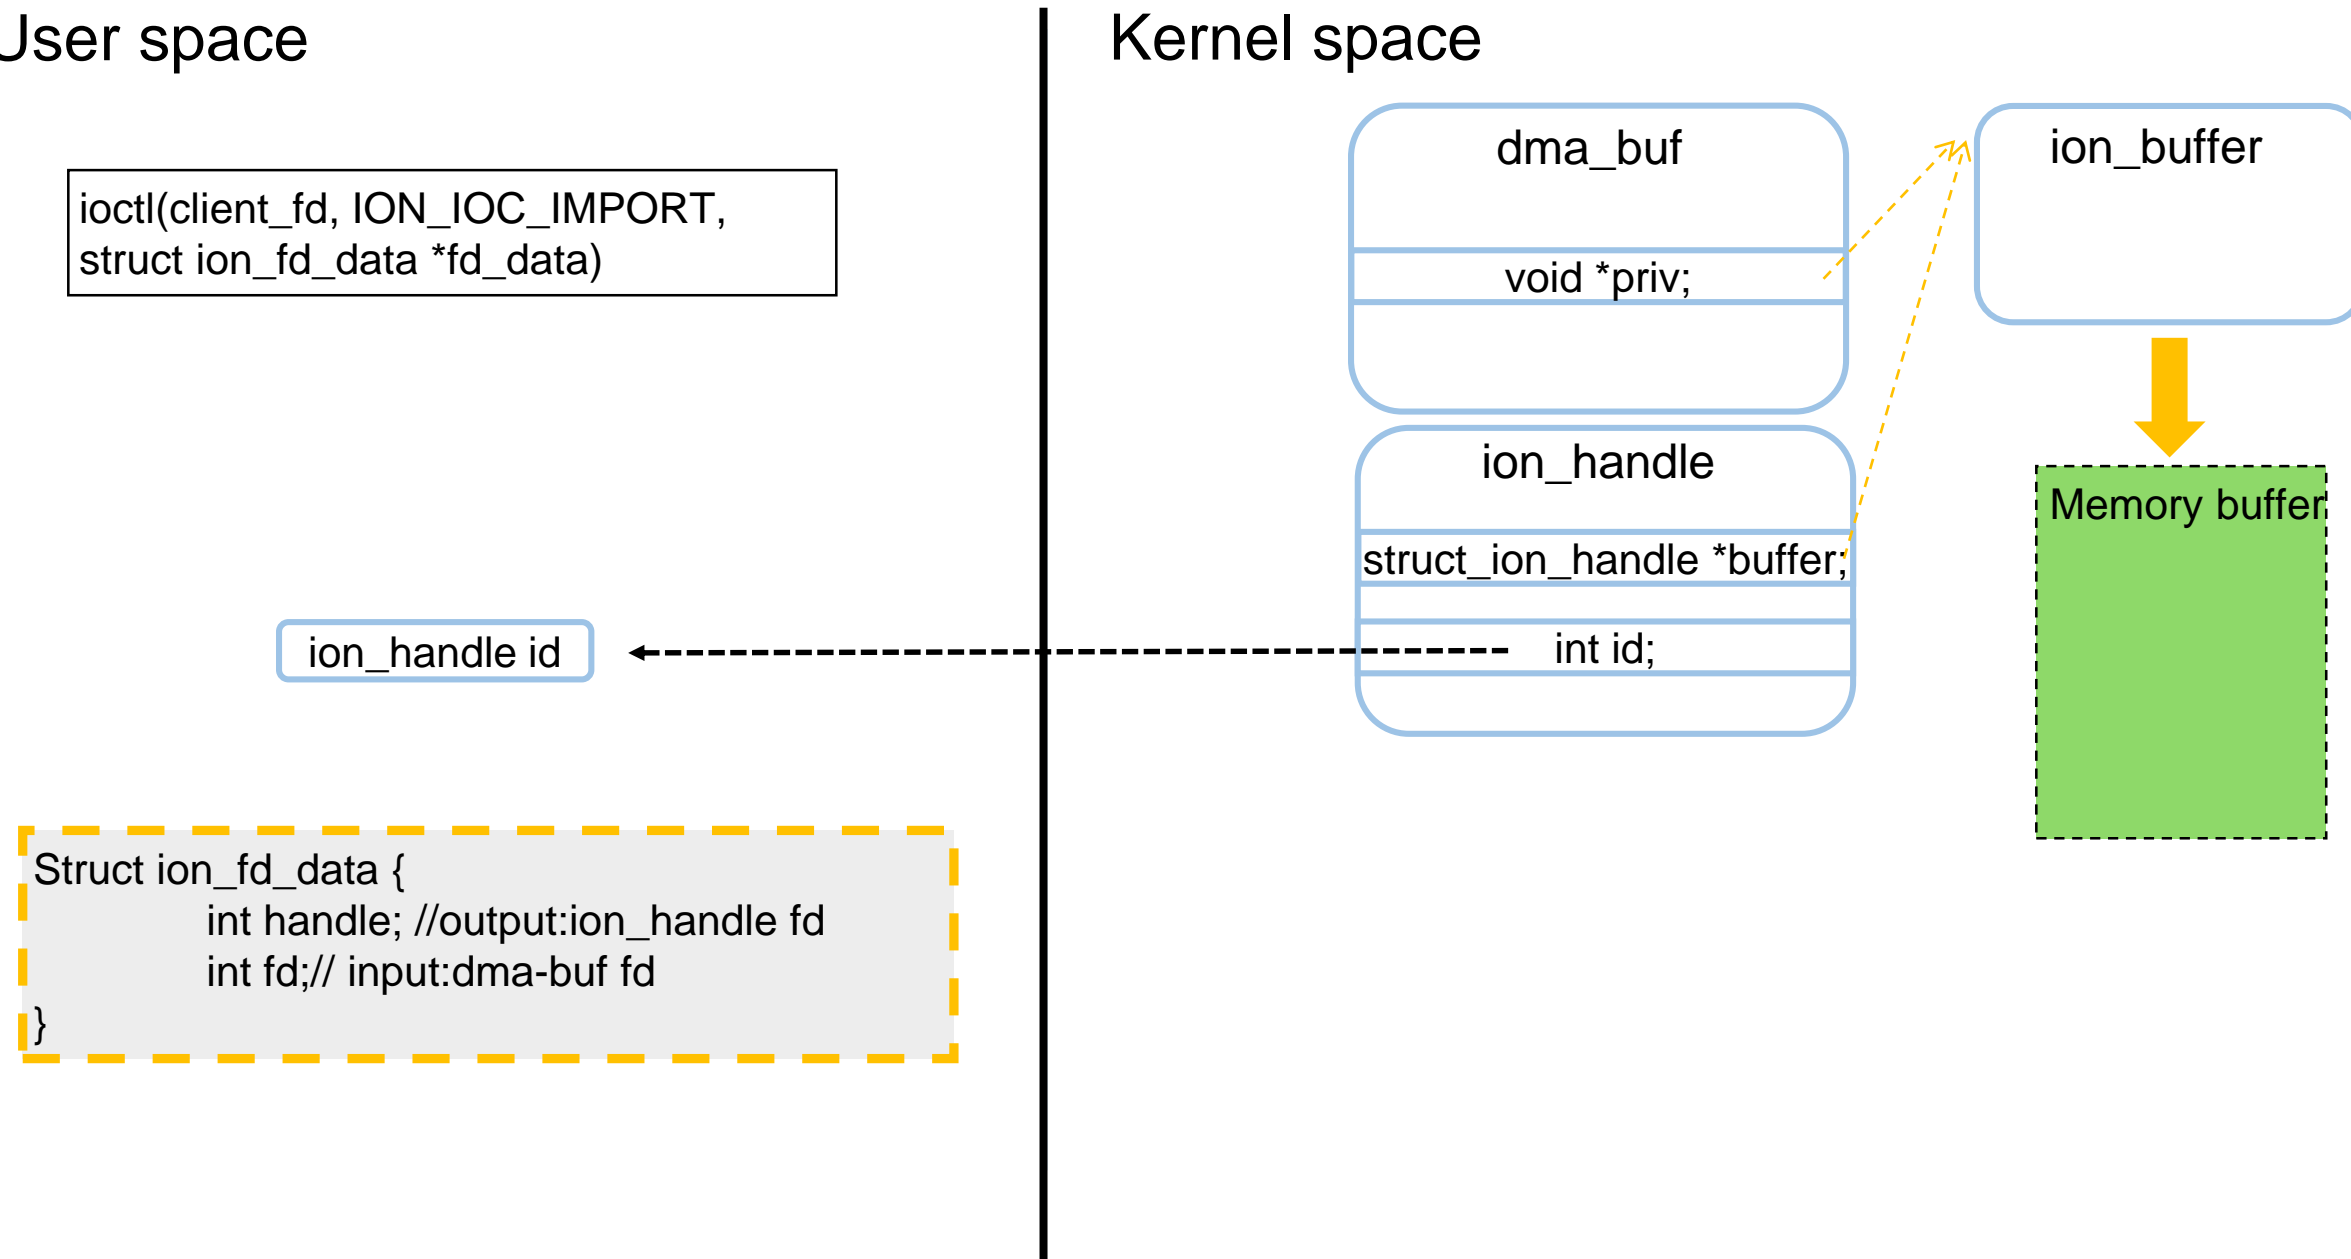
ion\_handle

struct\_ion\_handle \*buffer;

int id;

ion\_buffer

Memory buffer



# Using ION

- Allocate memory buffer from kernel space

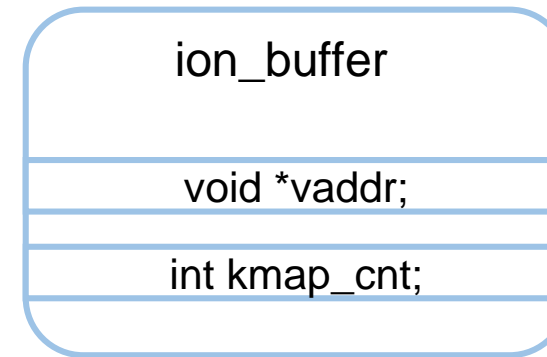
```
struct ion_handle *ion_alloc(struct ion_client *client, size_t len,  
                             size_t align, unsigned int heap_id_mask,  
                             unsigned int flags);
```

- Free memory buffer from kernel space

```
void ion_free(struct ion_client *client,  
              struct ion_handle *handle);
```

- Map the memory buffer into kernel space

```
void *ion_map_kernel(struct ion_client *client, struct ion_handle *handle);
```



# Diving into ExplosION

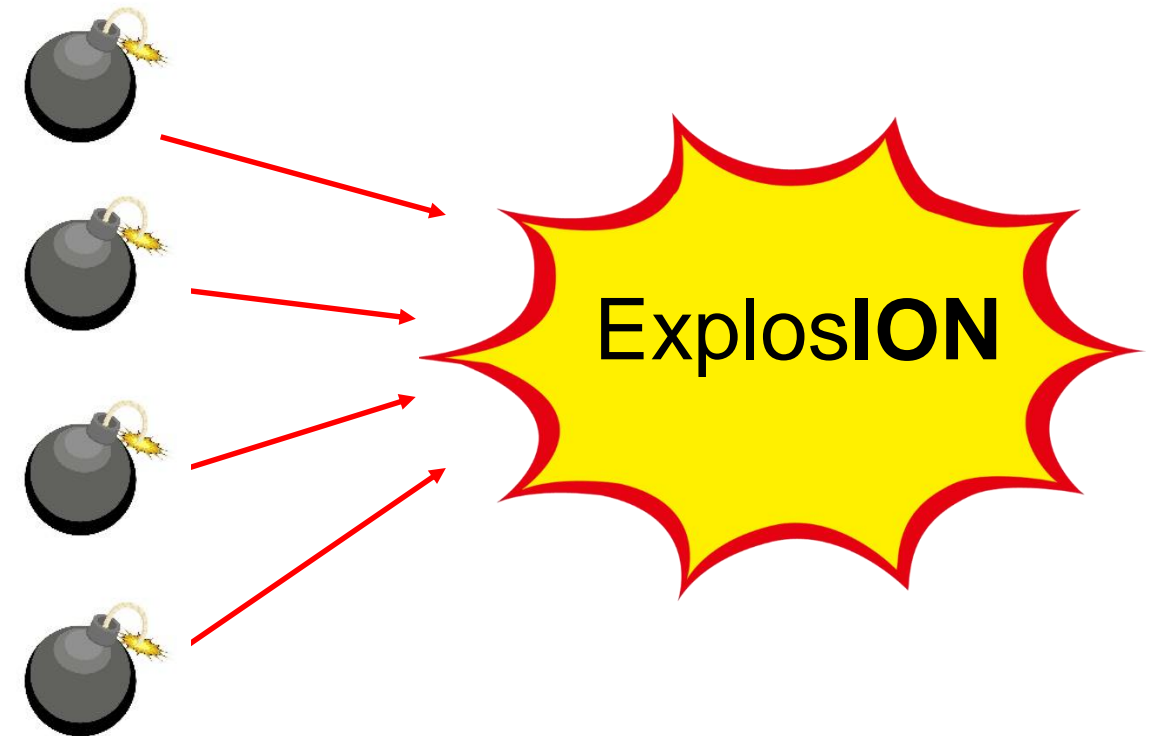
The Characteristics of ION:

A base driver used by vendor drivers

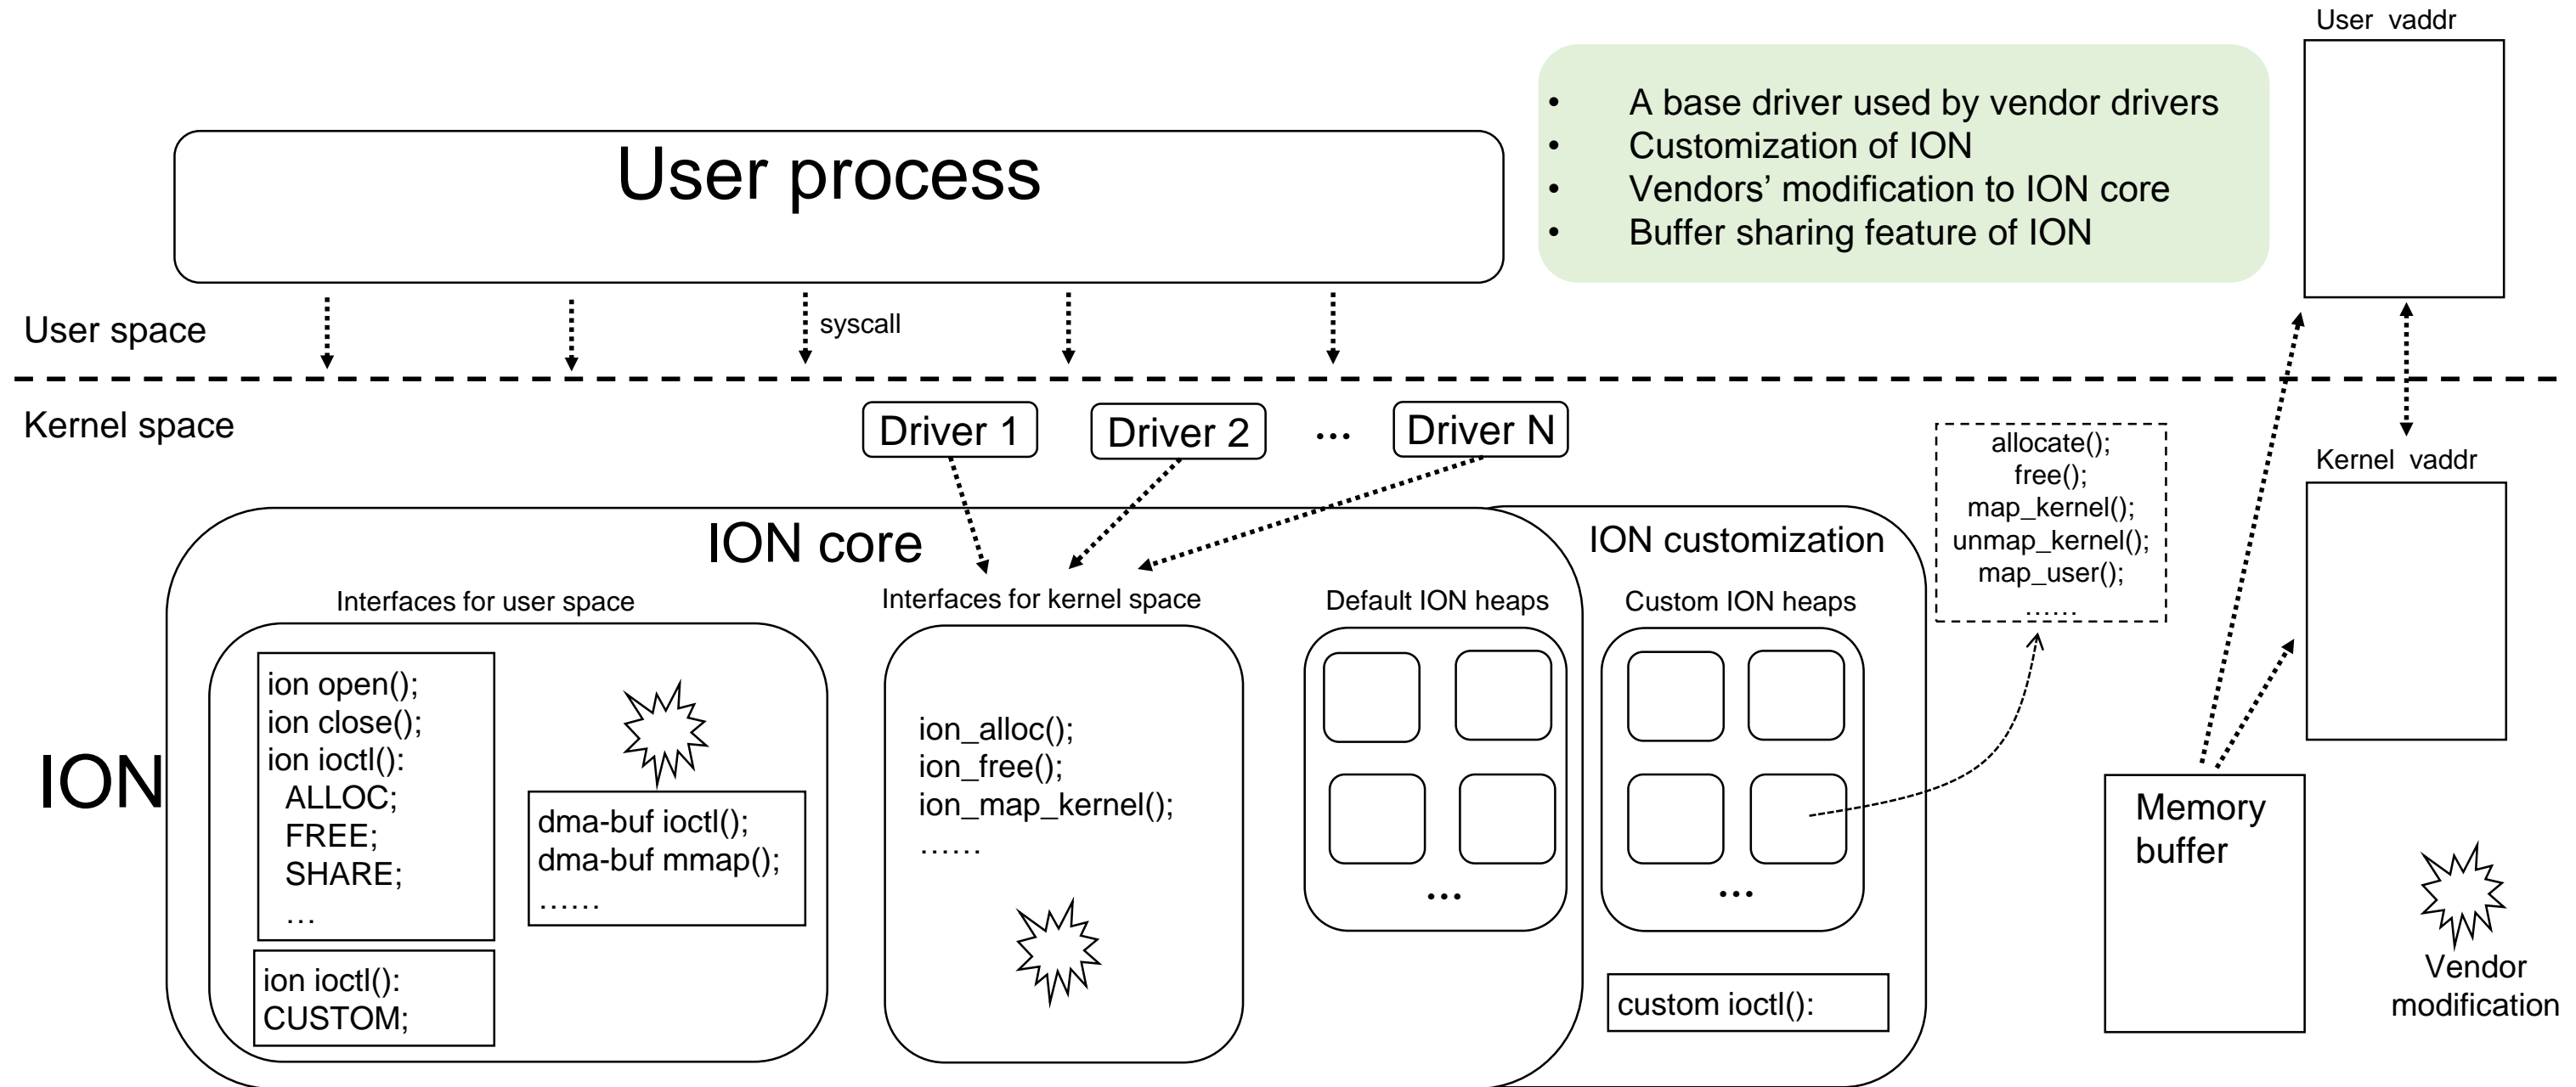
Customization of ION

Vendors' modification to ION core

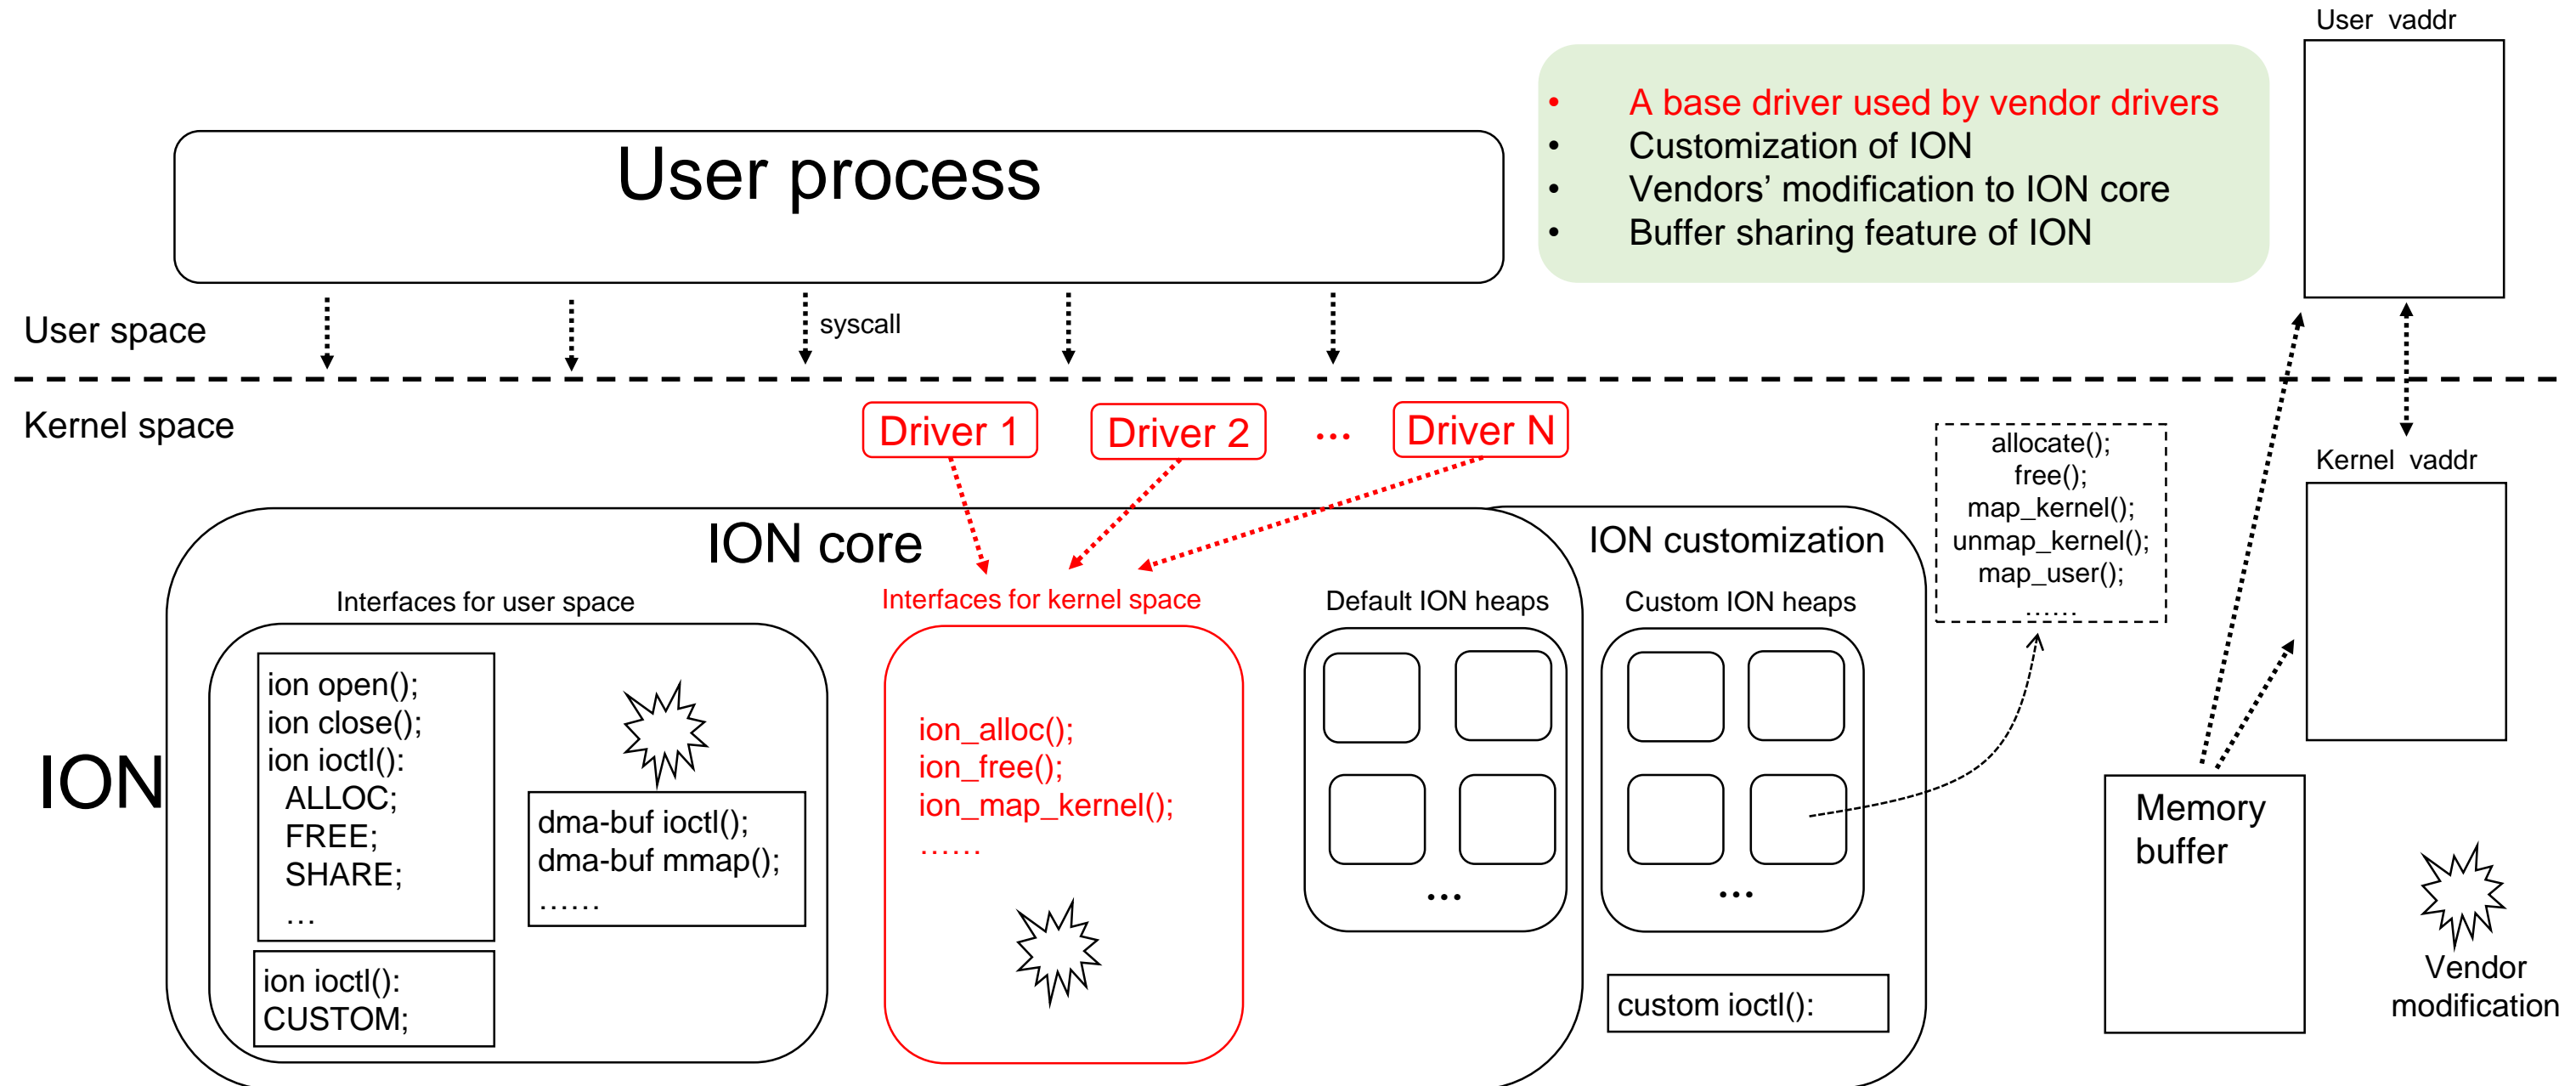
Buffer sharing feature of ION



# Legacy ION Architecture

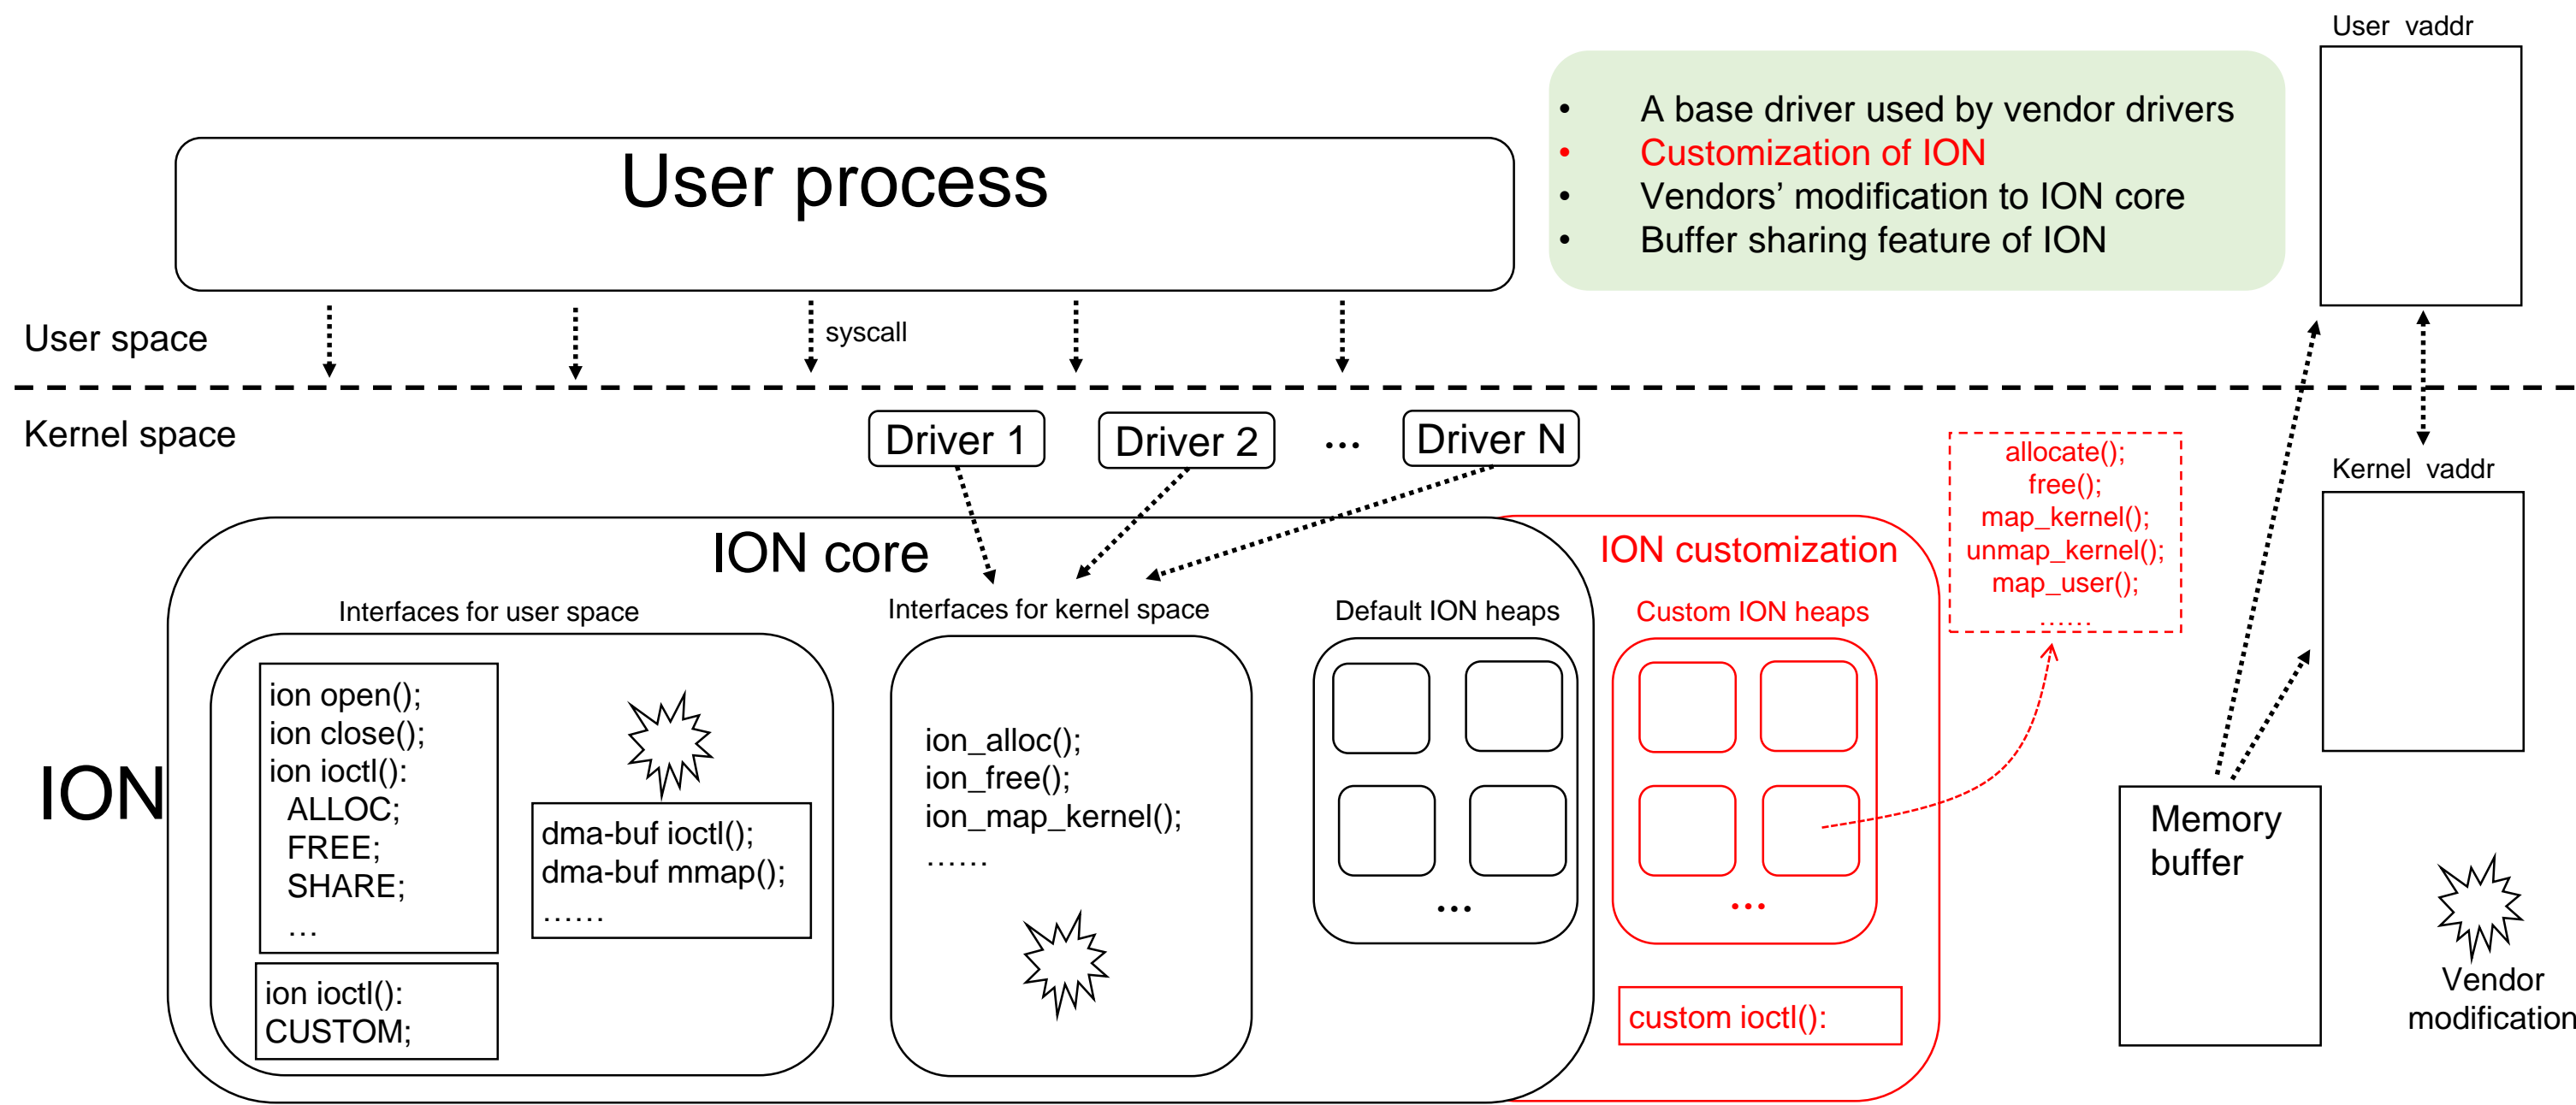


# Legacy ION Architecture

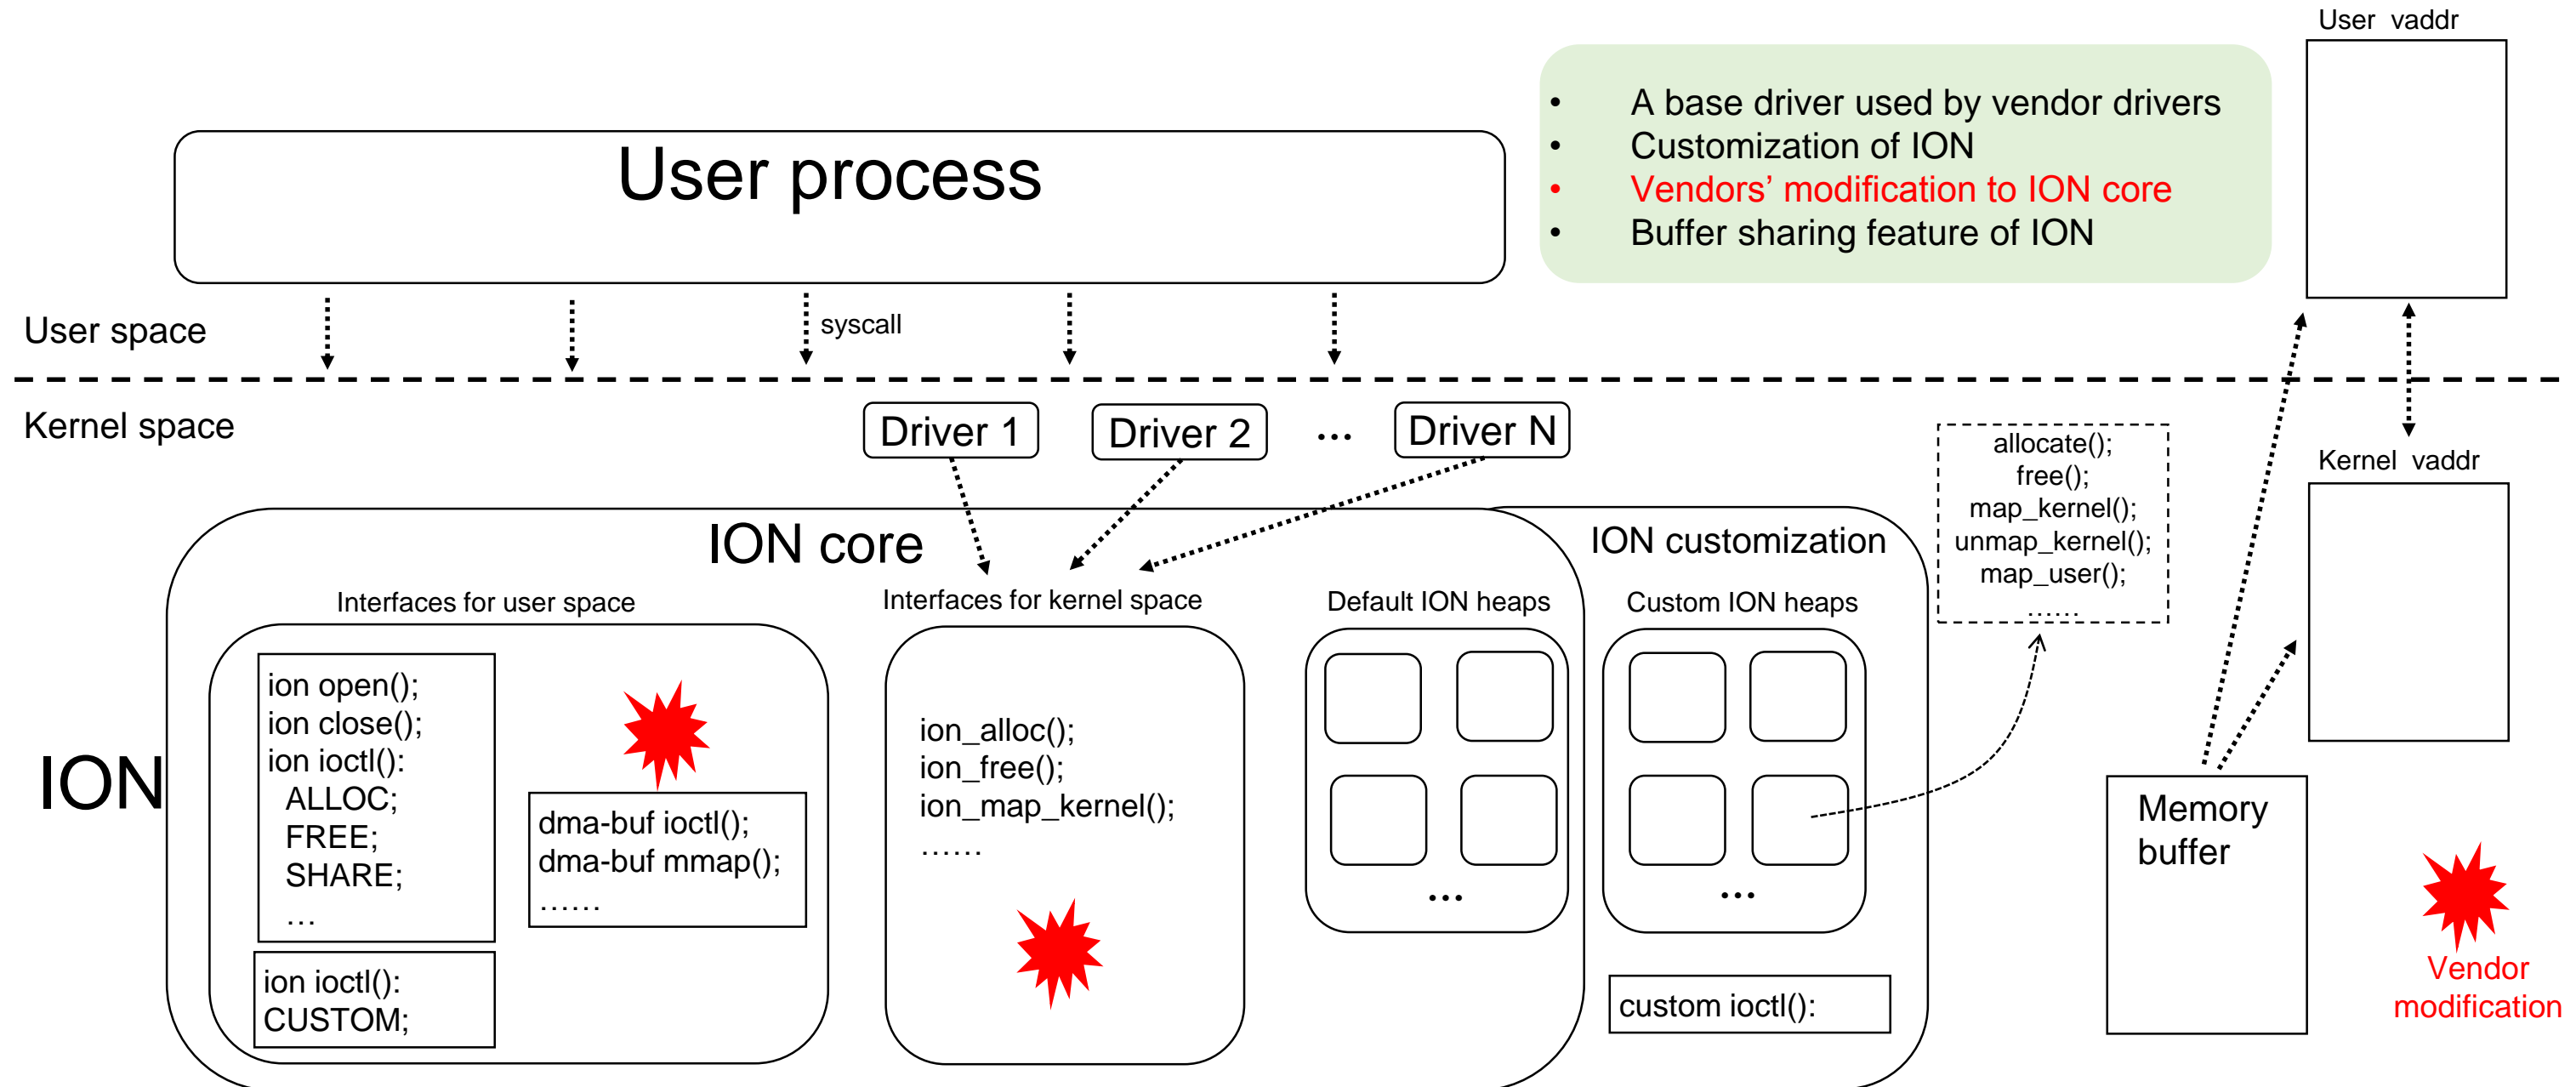




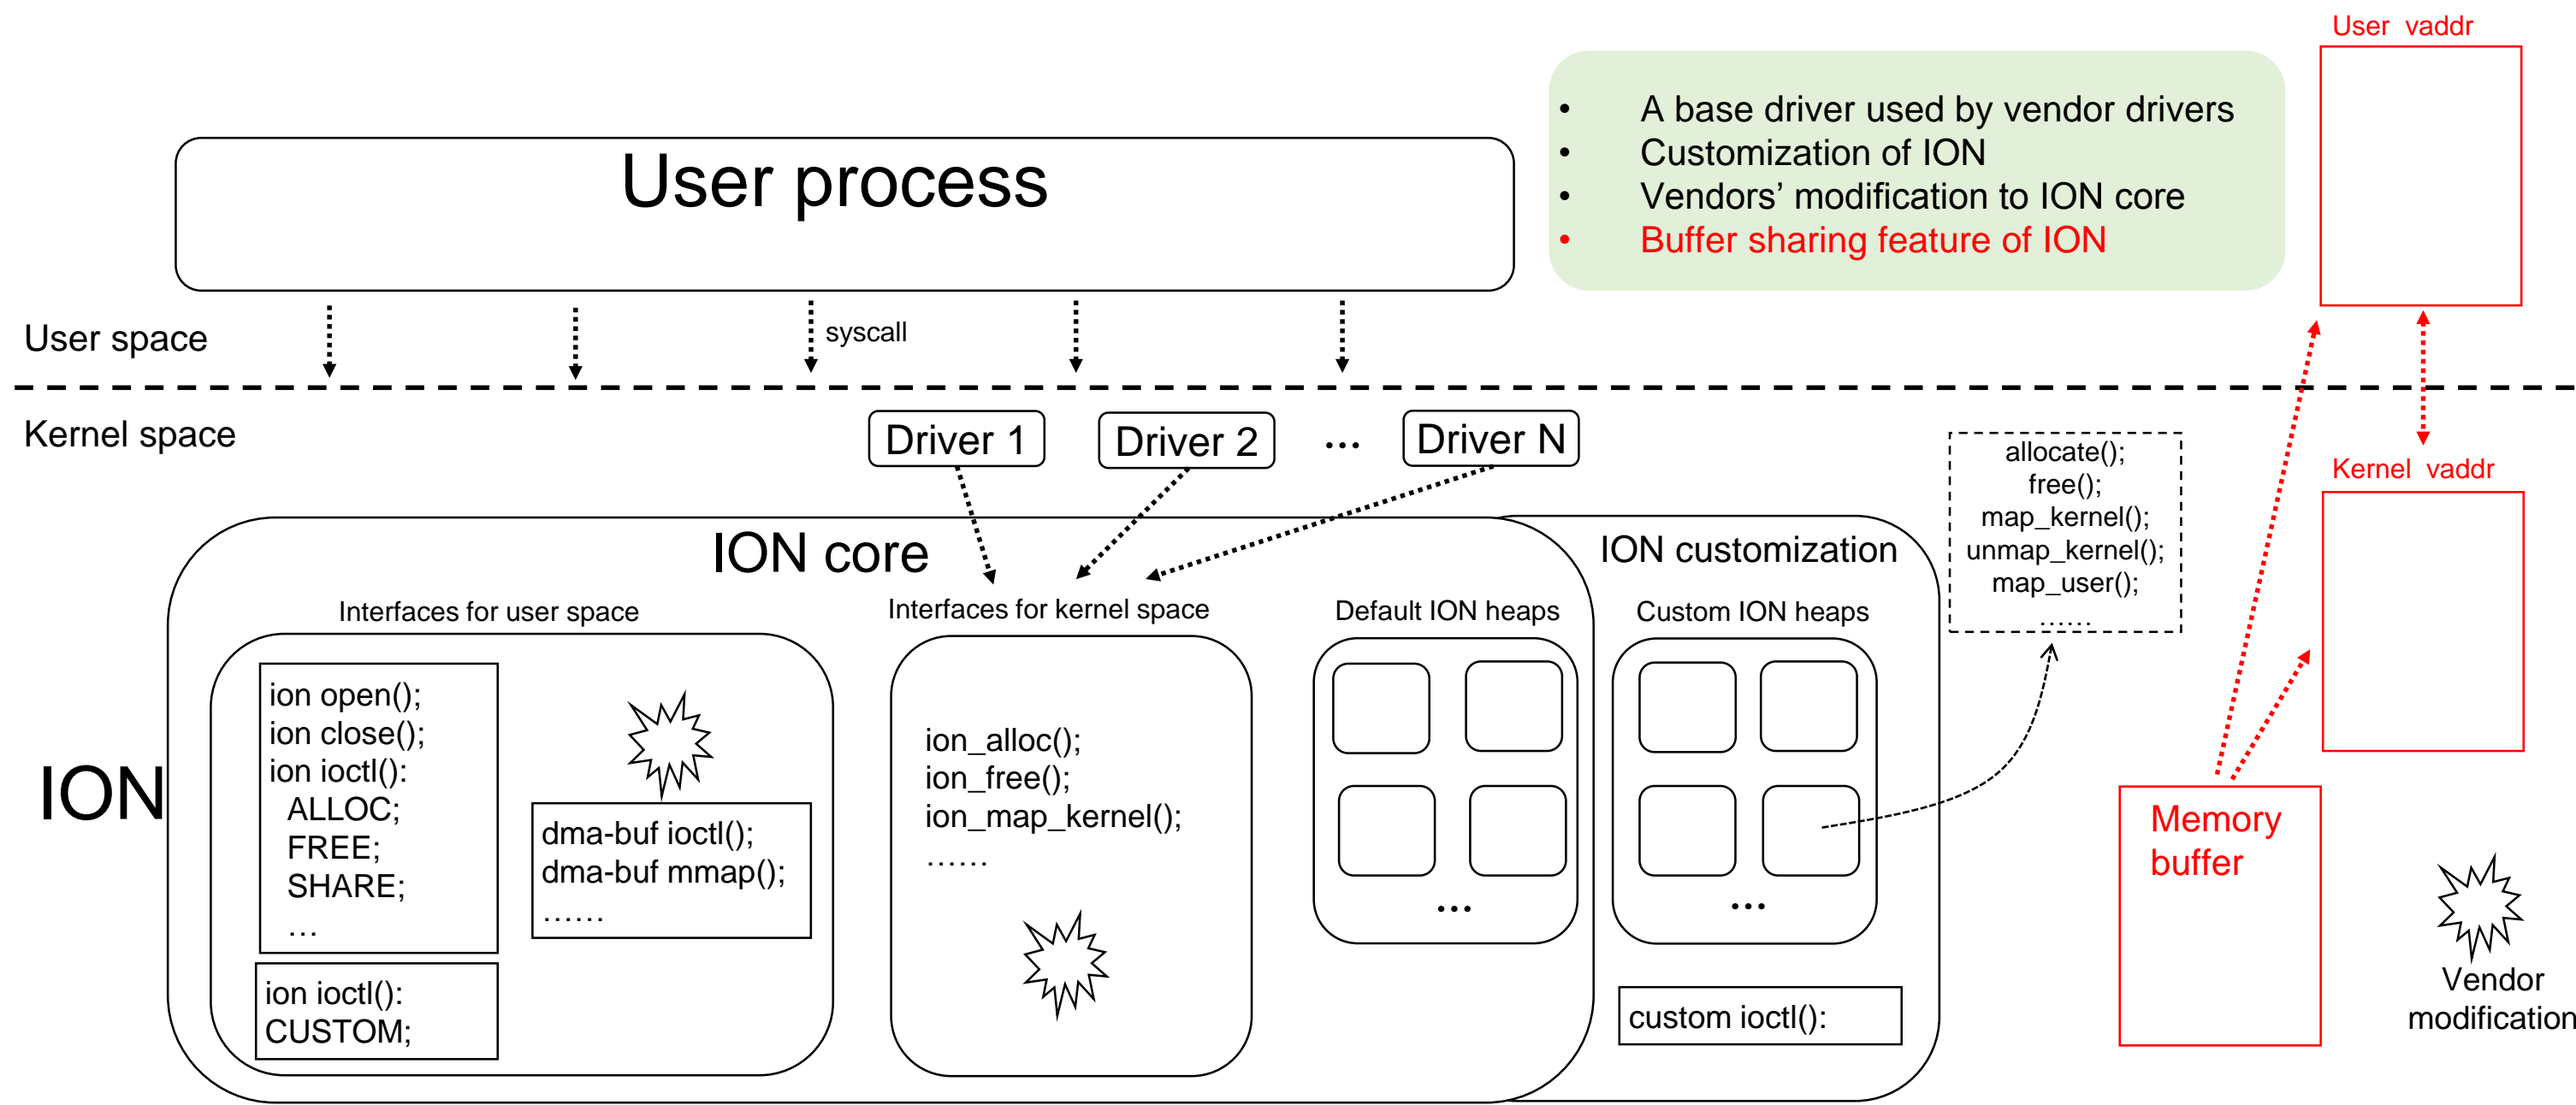
# Legacy ION Architecture



# Legacy ION Architecture

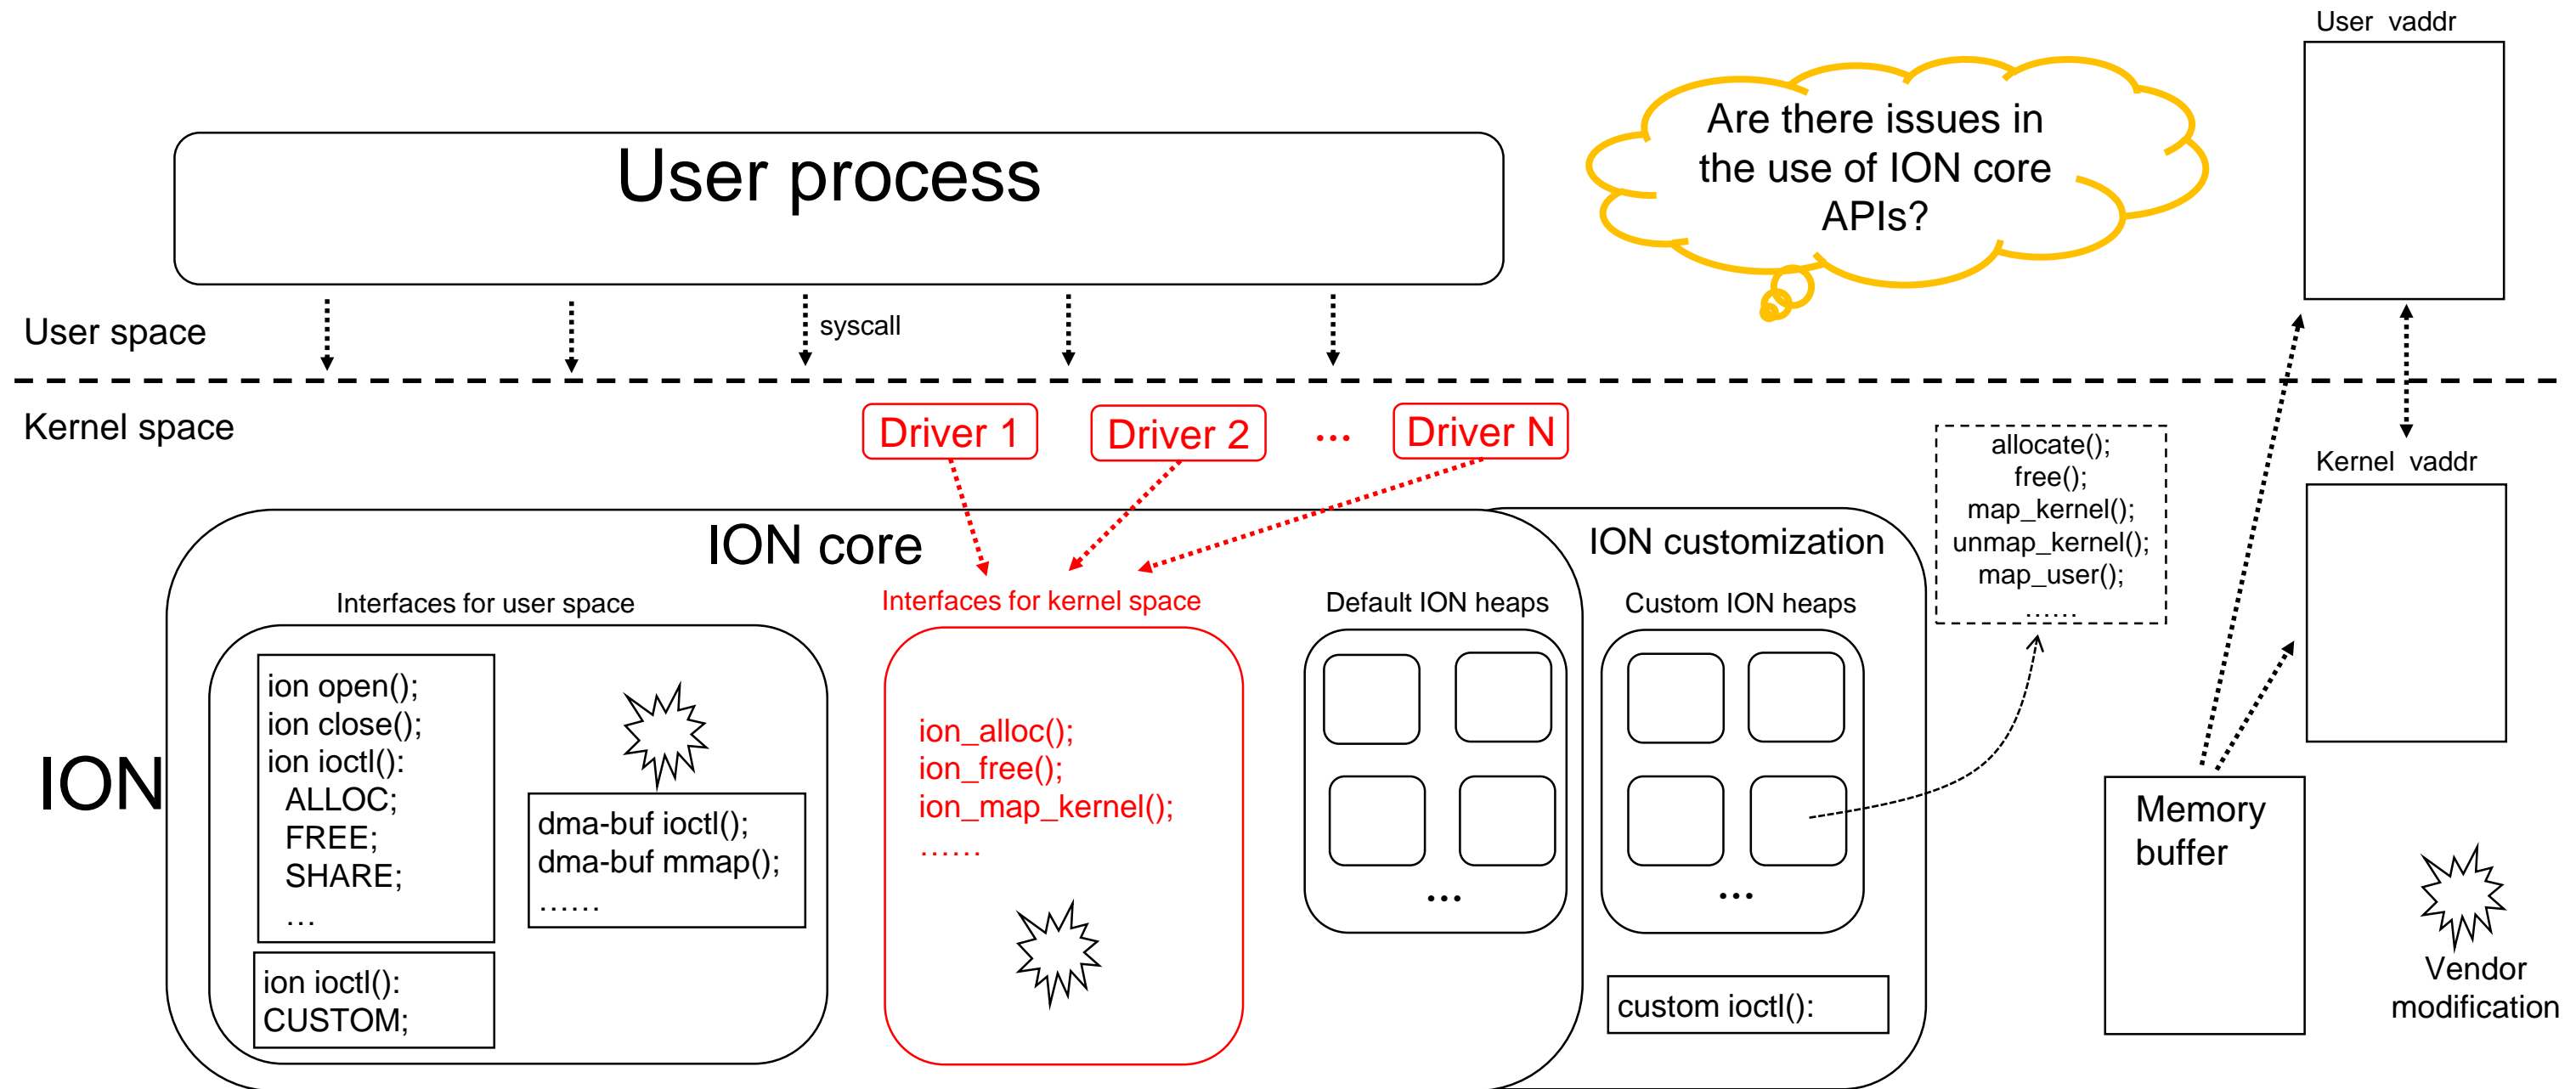


# Legacy ION Architecture



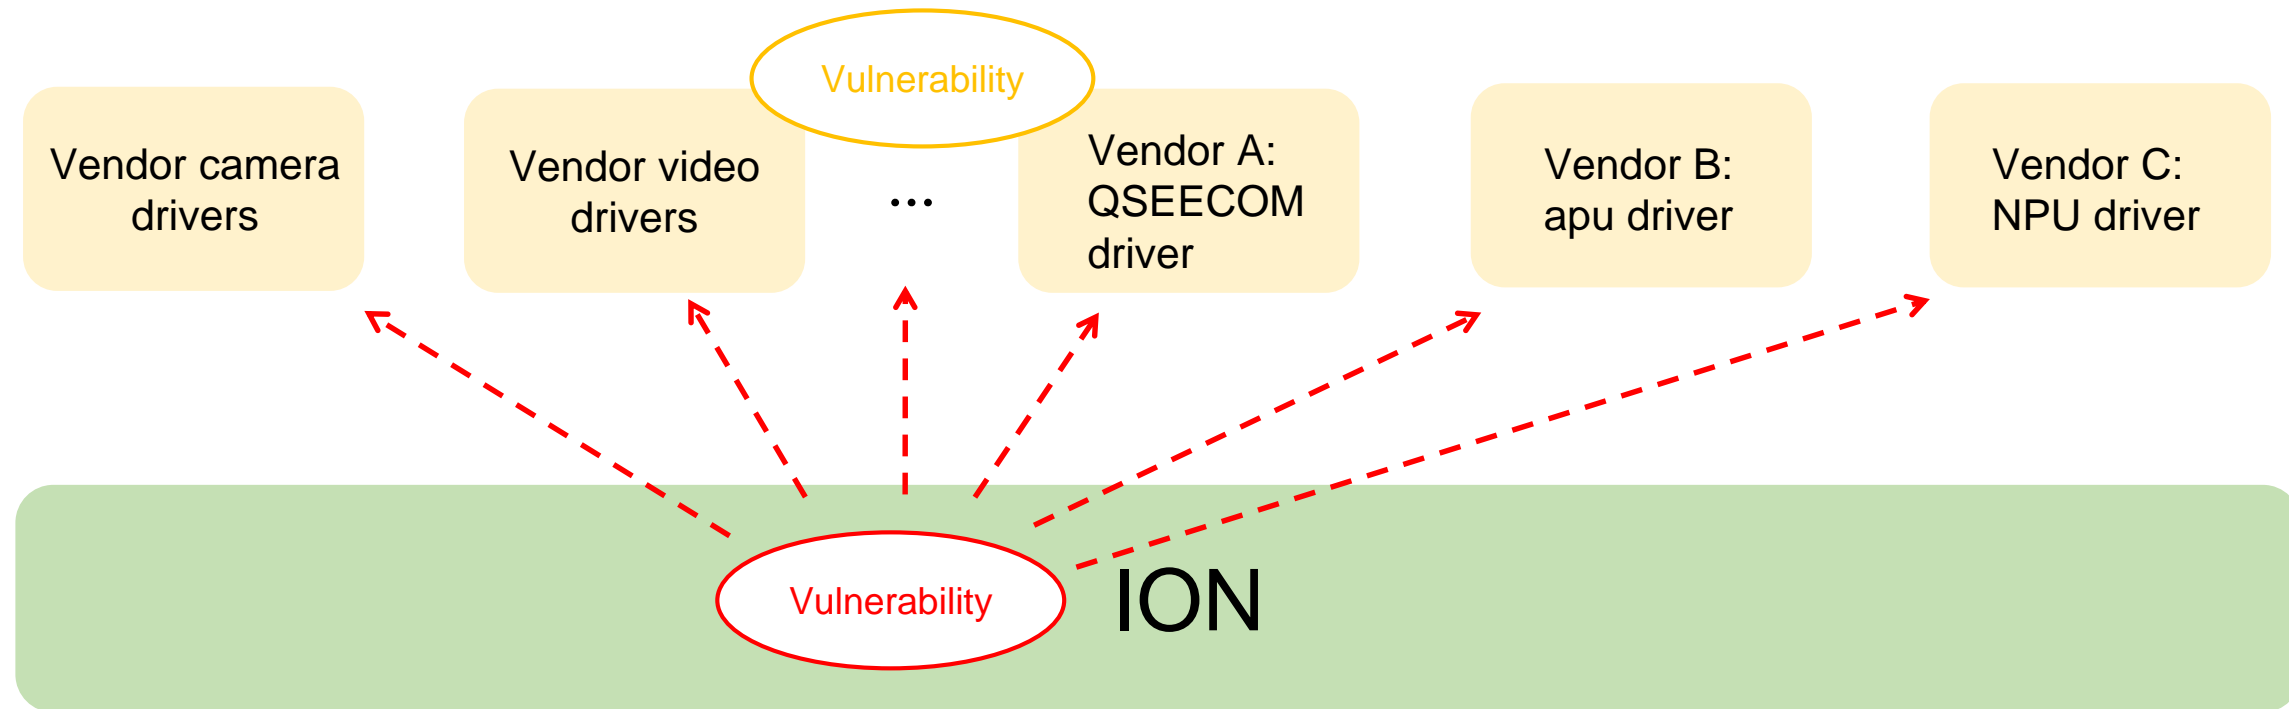


# A base driver used by vendor drivers





## A base driver used by vendor drivers




Vulnerability: From ION Core



Vulnerability: Misuse of ION Core APIs



## A base driver used by vendor drivers

- Analyze the design and implementation of ION core APIs for kernel space
- Check if vendor drivers use ION core APIs correctly



- Vulnerabilities from ION core:  
CVE-2021-0929   
CVE-2021-39714
- Misuse of ION Core APIs by vendor driver:  
CVE-2022-20110 

# CVE-2021-0929

- Affected ION version:  
ION V1,V2,V3
- Root cause:  
ION core exposed an API to user space which it shouldn't be
- Impact:  
An UAF issue would be introduced into vendor drivers by ION core in a specific use scenario

# CVE-2021-0929

## An ION use scenario

### User Process

Create dma-buf fd with ION

dma-buf fd

### Vendor driver

Import dma-buf fd, and create an ion\_handle:  
`ion_handle = ion_import_dma_buf_fd(ion_client, dma-buf fd);`

Map the memory buffer represented by ion\_handle into kernel space:

`kernel_vaddr = ion_map_kernel(ion_client, ion_handle);`  
(the kernel vaddr will be saved into `ion_buffer->vaddr`)

Reference the `kernel_vaddr`;



# CVE-2021-0929

```
static long dma_buf_ioctl(struct file *file,
                          unsigned int cmd, unsigned long arg)
{
    .....
    dmabuf = file->private_data;
    switch (cmd) {
    case DMA_BUF_IOCTL_SYNC:
        .....
        if (sync.flags & DMA_BUF_SYNC_END)
            ret = dma_buf_end_cpu_access(dmabuf,
            direction);
        .....
    }
```

```
int dma_buf_end_cpu_access(struct dma_buf *dmabuf,
                           enum dma_data_direction direction)
{
    int ret = 0;
    WARN_ON(!dmabuf);
    if (dmabuf->ops->end_cpu_access)
        ret = dmabuf->ops->end_cpu_access(dmabuf, direction);
    return ret;
}
```

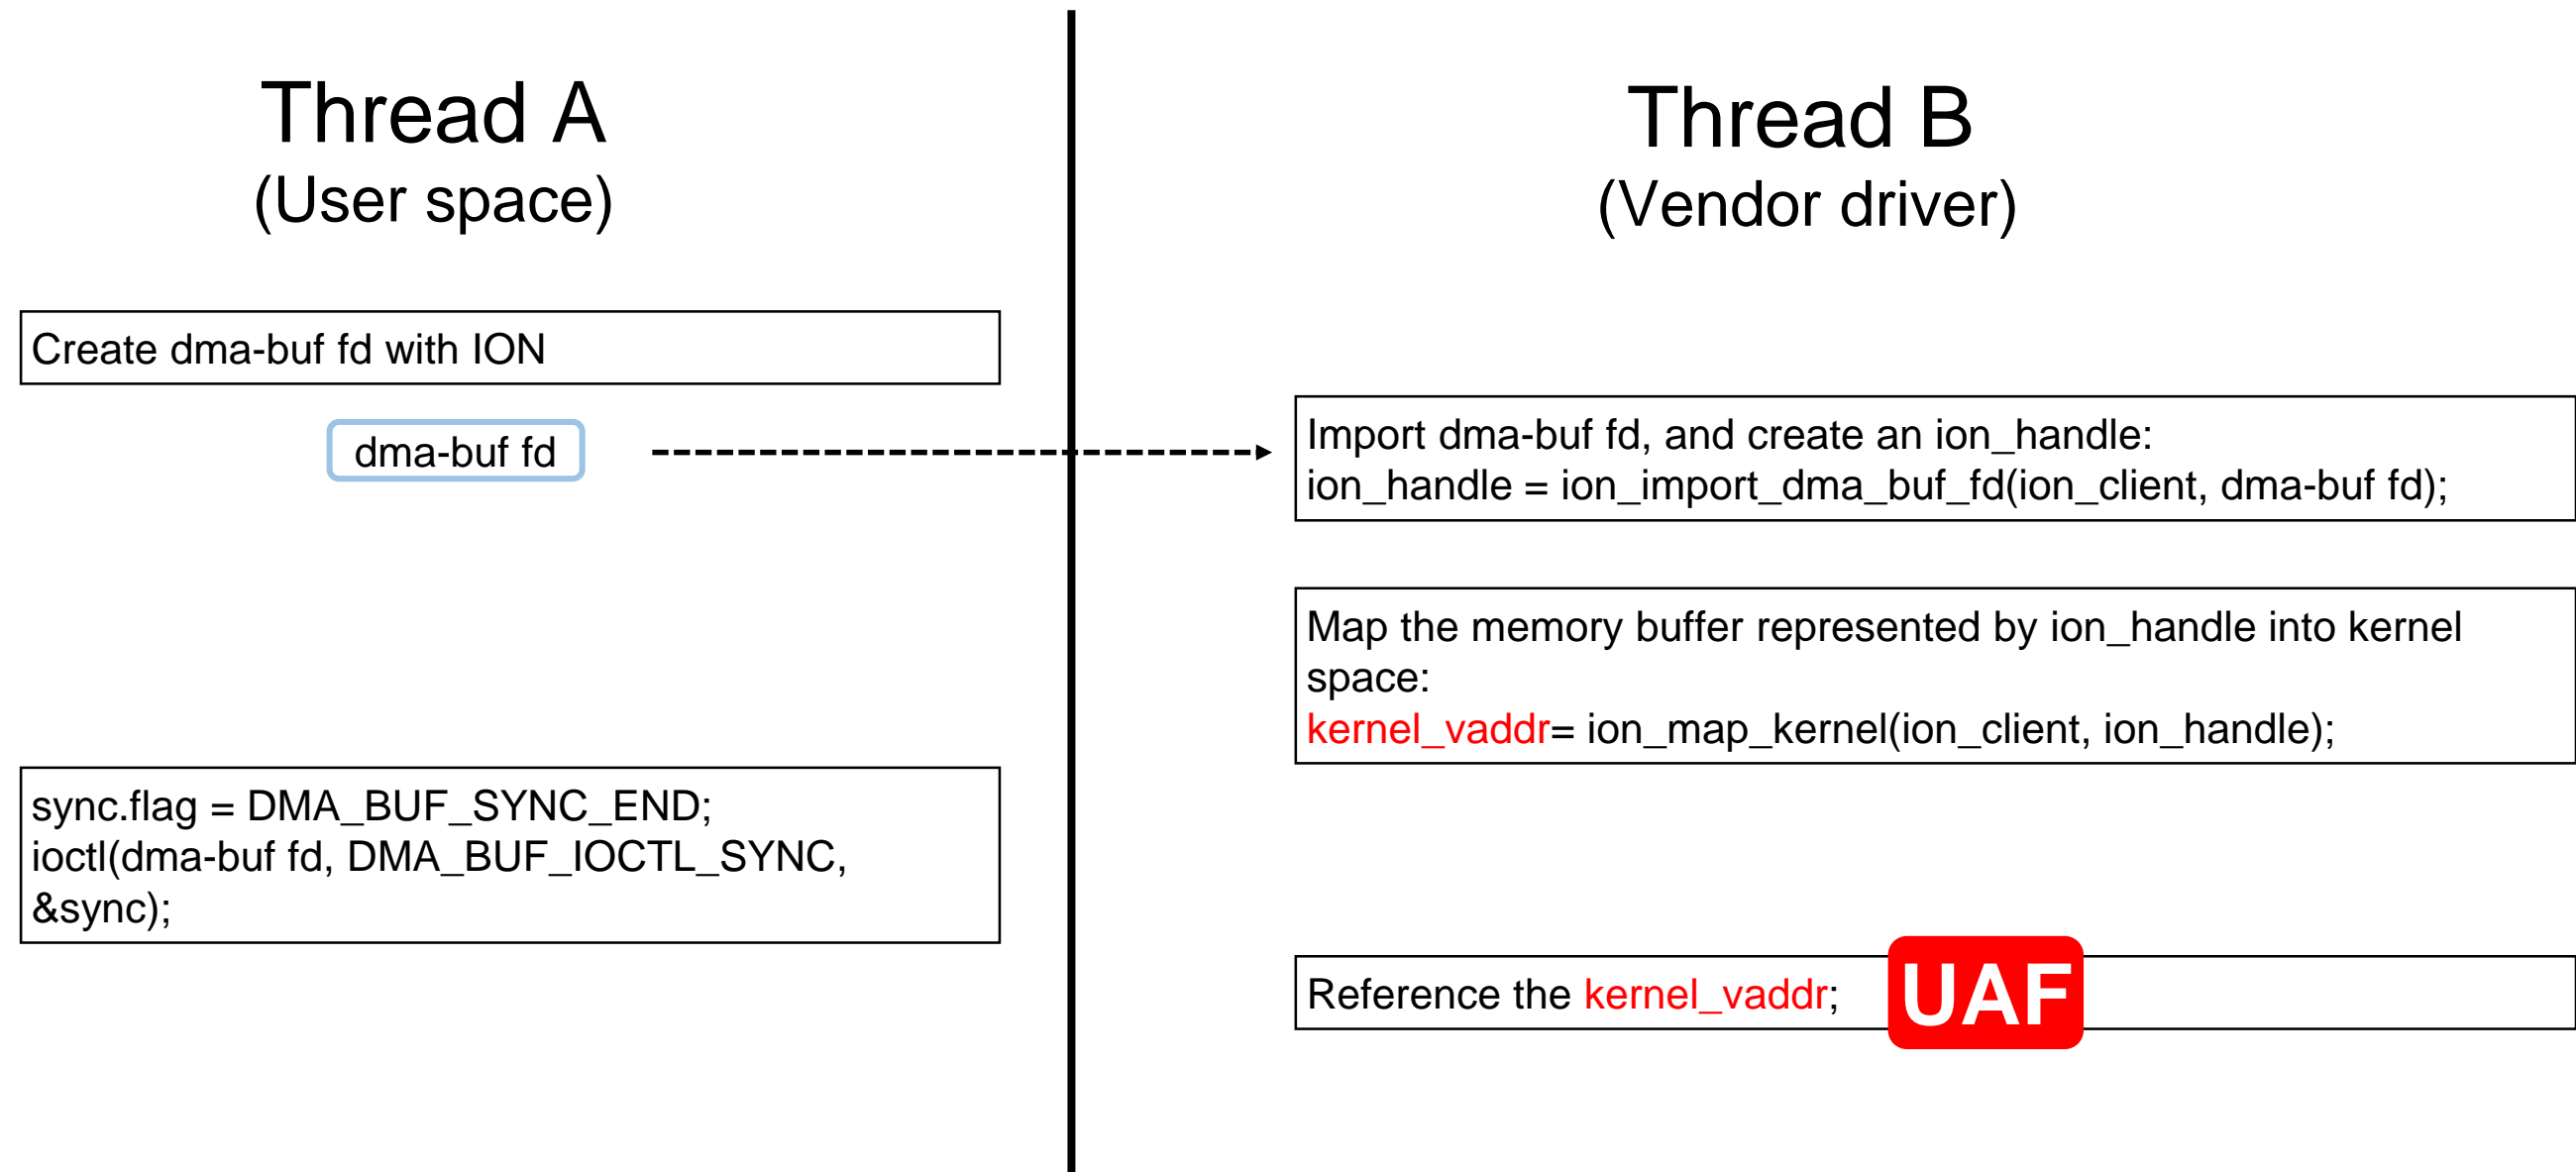
```
static int ion_dma_buf_end_cpu_access(struct dma_buf *dmabuf,
                                      ...)
{
    struct ion_buffer *buffer = dmabuf->priv;
    mutex_lock(&buffer->lock);
    ion_buffer_kmap_put(buffer);
    .....
}
```

```
void ion_heap_unmap_kernel(struct ion_heap *heap,
                           struct ion_buffer *buffer)
{
    .....
    vunmap(buffer->vaddr);
}
```

buffer->vaddr gets  
unmapped in  
kernel space!!!

# CVE-2021-0929

UAF would happen in a race condition:



# CVE-2022-20110

- Affected ION version:  
A vendor's devices with ION V1
- Root cause:  
Misuse of ION core APIs by vendor driver
- Impact:  
UAF

# CVE-2022-20110

```
void *ion_map_kernel(struct ion_client *client, struct
ion_handle *handle)
{
    struct ion_buffer *buffer;
    void *vaddr;

    mutex_lock(&client->lock);
    .....
    buffer = handle->buffer;
    .....
    mutex_lock(&buffer->lock);
    vaddr = ion_handle_kmap_get(handle);
    mutex_unlock(&buffer->lock);
    mutex_unlock(&client->lock);
    return vaddr;
}
```

```
static void *ion_buffer_kmap_get(struct ion_buffer *buffer)
{
    void *vaddr;

    if (buffer->kmap_cnt) {
        buffer->kmap_cnt++;
        return buffer->vaddr;
    }
    vaddr = buffer->heap->ops->map_kernel(buffer->heap,
buffer);
    .....
    buffer->vaddr = vaddr;
    buffer->kmap_cnt++;
    return vaddr;
}
```

buffer->vaddr and buffer->kmap\_cnt are protected by mutex locks!

# CVE-2022-20110

Access internal data instead of using the exported API:

```
static long ion_sys_cache_sync(struct ion_client *client,
                              struct ion_sys_cache_sync_param *param,
                              int from_kernel)
{
    .....
    unsigned long sync_va = 0;
    struct ion_buffer *buffer;

    .....
    if (buffer->kmap_cnt != 0) {
        sync_va = (unsigned long)buffer->vaddr;
    } else {
        sync_va = (unsigned long)ion_map_kernel(client, kernel_handle);
        ion_need_unmap_flag = 1;
    }

    .....
    ret = __cache_sync_by_range(client, sync_type,
                                sync_va, ...);
    .....
    ion_unmap_kernel(client, kernel_handle);
    .....
}
```

buffer->vaddr and buffer->kmap\_cnt are not protected by mutex locks!

**UAF**

# CVE-2022-20110

UAF would happen in a race condition:

## Thread A

(enter ion\_sys\_cache\_sync)

```
if (buffer->kmap_cnt != 0) {  
    sync_va = buffer->vaddr;  
} else {  
    .....  
}
```

```
__cache_sync_by_range(client, sync_type,  
                        sync_va, sync_size,  
from_kernel);
```

**UAF**

## Thread B

(enter ion\_sys\_cache\_sync)

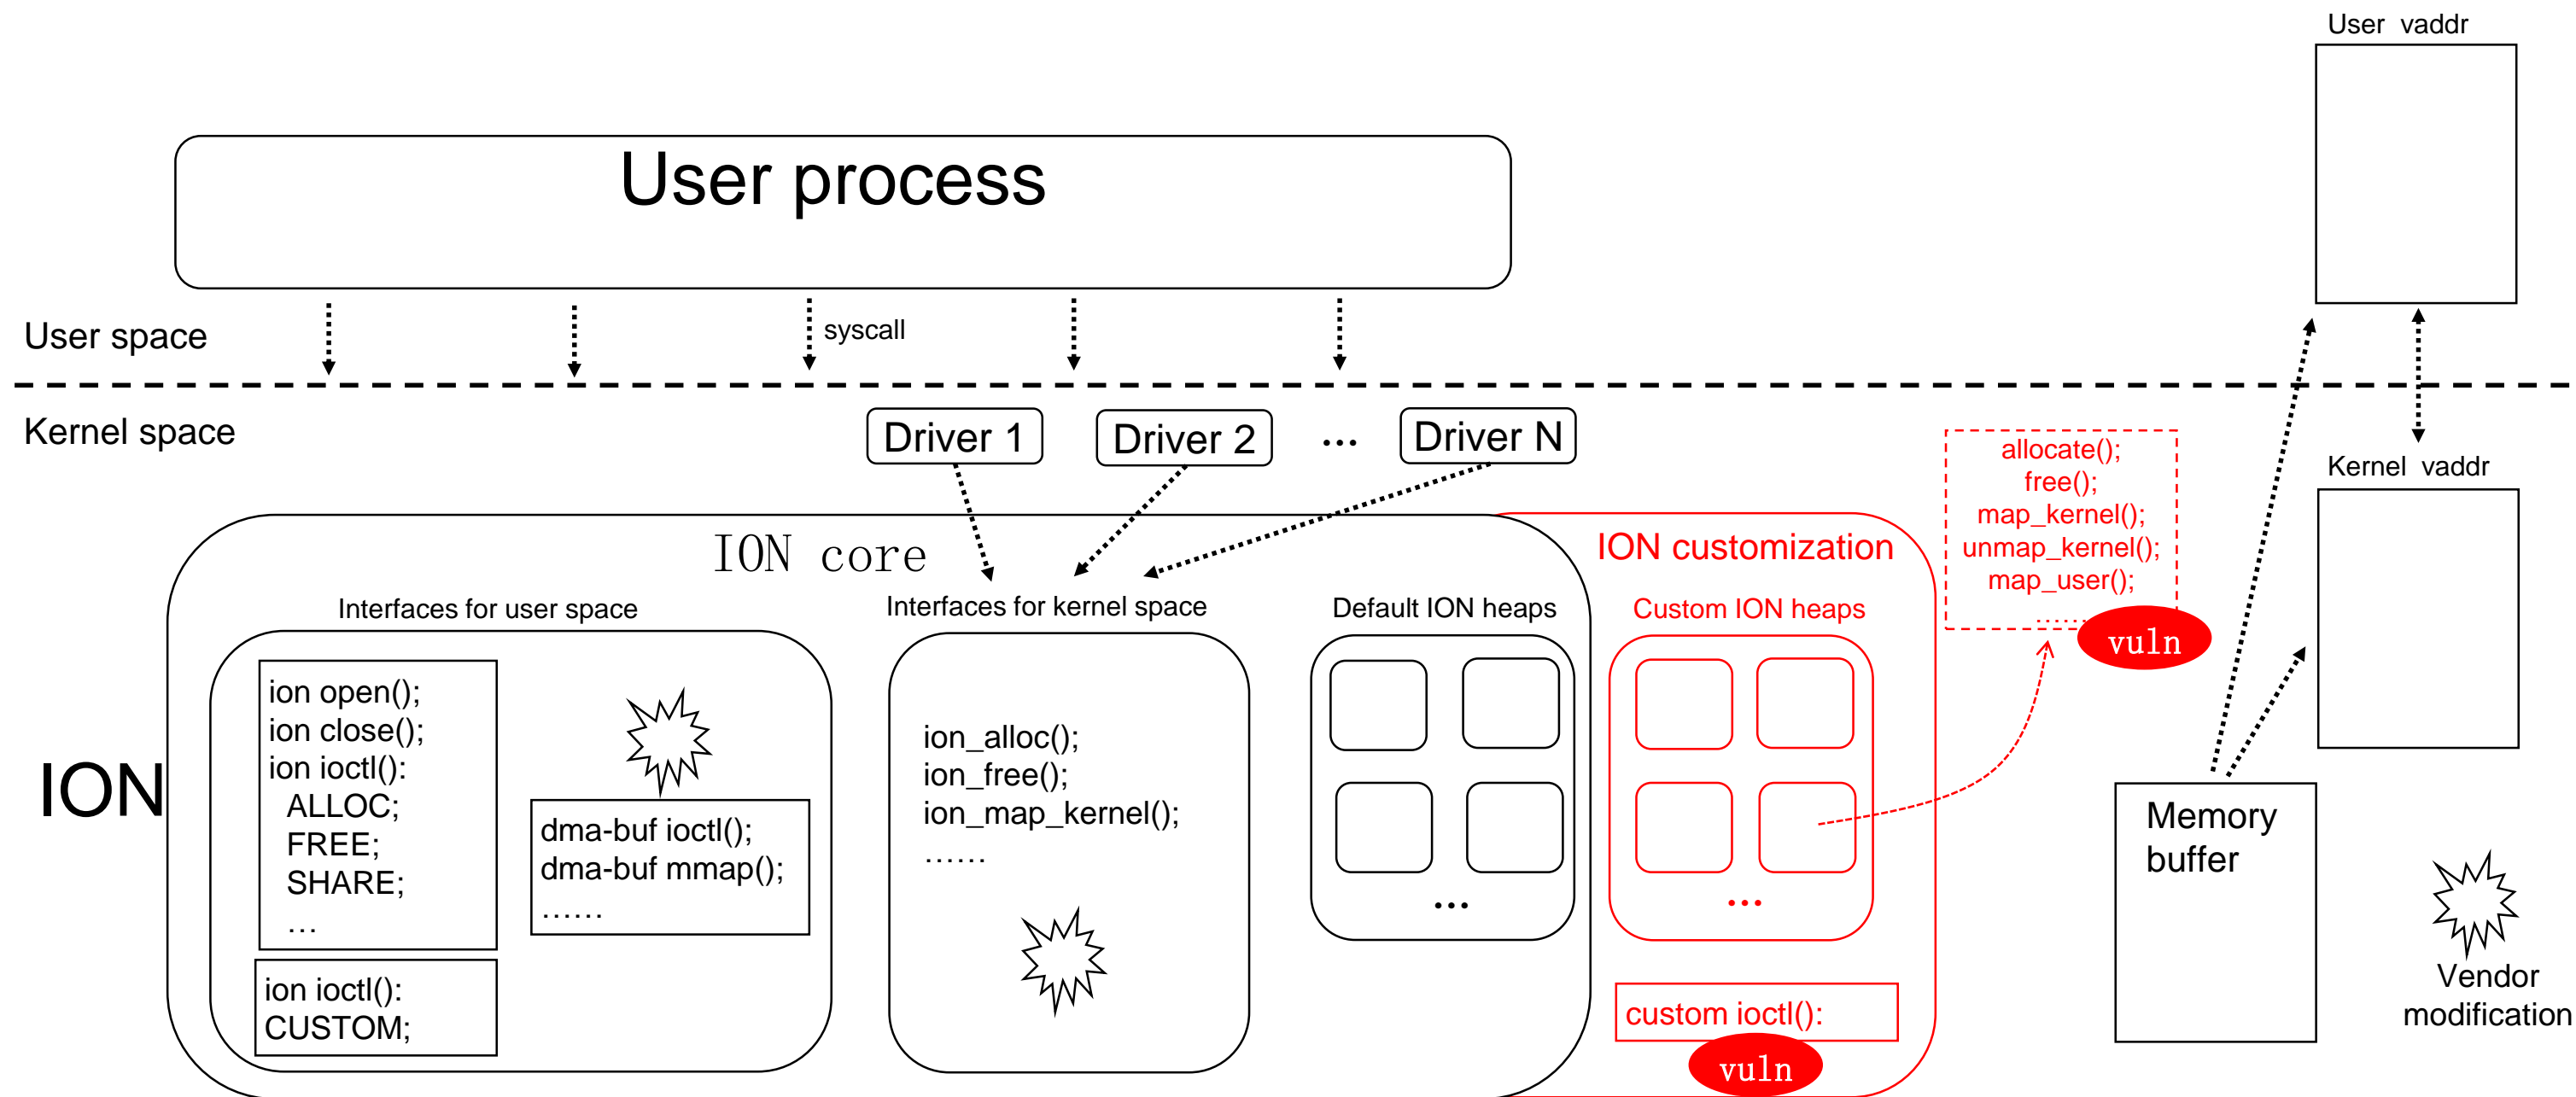
```
if (buffer->kmap_cnt != 0) {  
} else {  
    sync_va = ion_map_kernel(client, kernel_handle);  
}
```

```
__cache_sync_by_range(client, sync_type,  
                        sync_va, sync_size, from_kernel);
```

```
ion_unmap_kernel(client, kernel_handle);  
(sync_va will be unmapped !!!)
```



# Customization of ION







# Customization of ION

We found:

- ✓ Most vendors would add their own ion heaps
- ✓ Most vendors would add custom ioctl commands

Vendor A

```
long msm_ion_custom_ioctl(struct ion_client *client,
                          unsigned int cmd,
                          unsigned long arg)
{
    .....
    switch (cmd) {
    case ION_IOC_CLEAN_CACHES:
    case ION_IOC_INV_CACHES:
    case ION_IOC_CLEAN_INV_CACHES:
    {
        .....
    }
```

Vendor B

```
static long _ion_ioctl(struct ion_client *client,
                      unsigned int cmd, unsigned long arg, int
                      from_kernel)
{
    .....
    switch (cmd) {
    case ION_CMD_SYSTEM:
        .....
    case ION_CMD_MULTIMEDIA:
        .....
        break;
    }
    .....
}
```

Customization is  
proven to be  
vulnerable in many  
cases!

# Dozens of vulnerabilities in the customization

ion\_heap customization:

CVE-2021-0498,CVE-2021-0528,CVE-2021-0489,CVE-2021-0493,CVE-2021-0492,  
CVE-2021-0490,CVE-2021-0496,CVE-2021-0526,CVE-2021-0420,CVE-2021-0525,  
CVE-2021-0495,CVE-2021-0497,CVE-2021-0491,CVE-2021-0421,CVE-2021-0530,  
CVE-2021-0494,CVE-2021-0424,CVE-2021-0527 ...

ion custom ioctl:

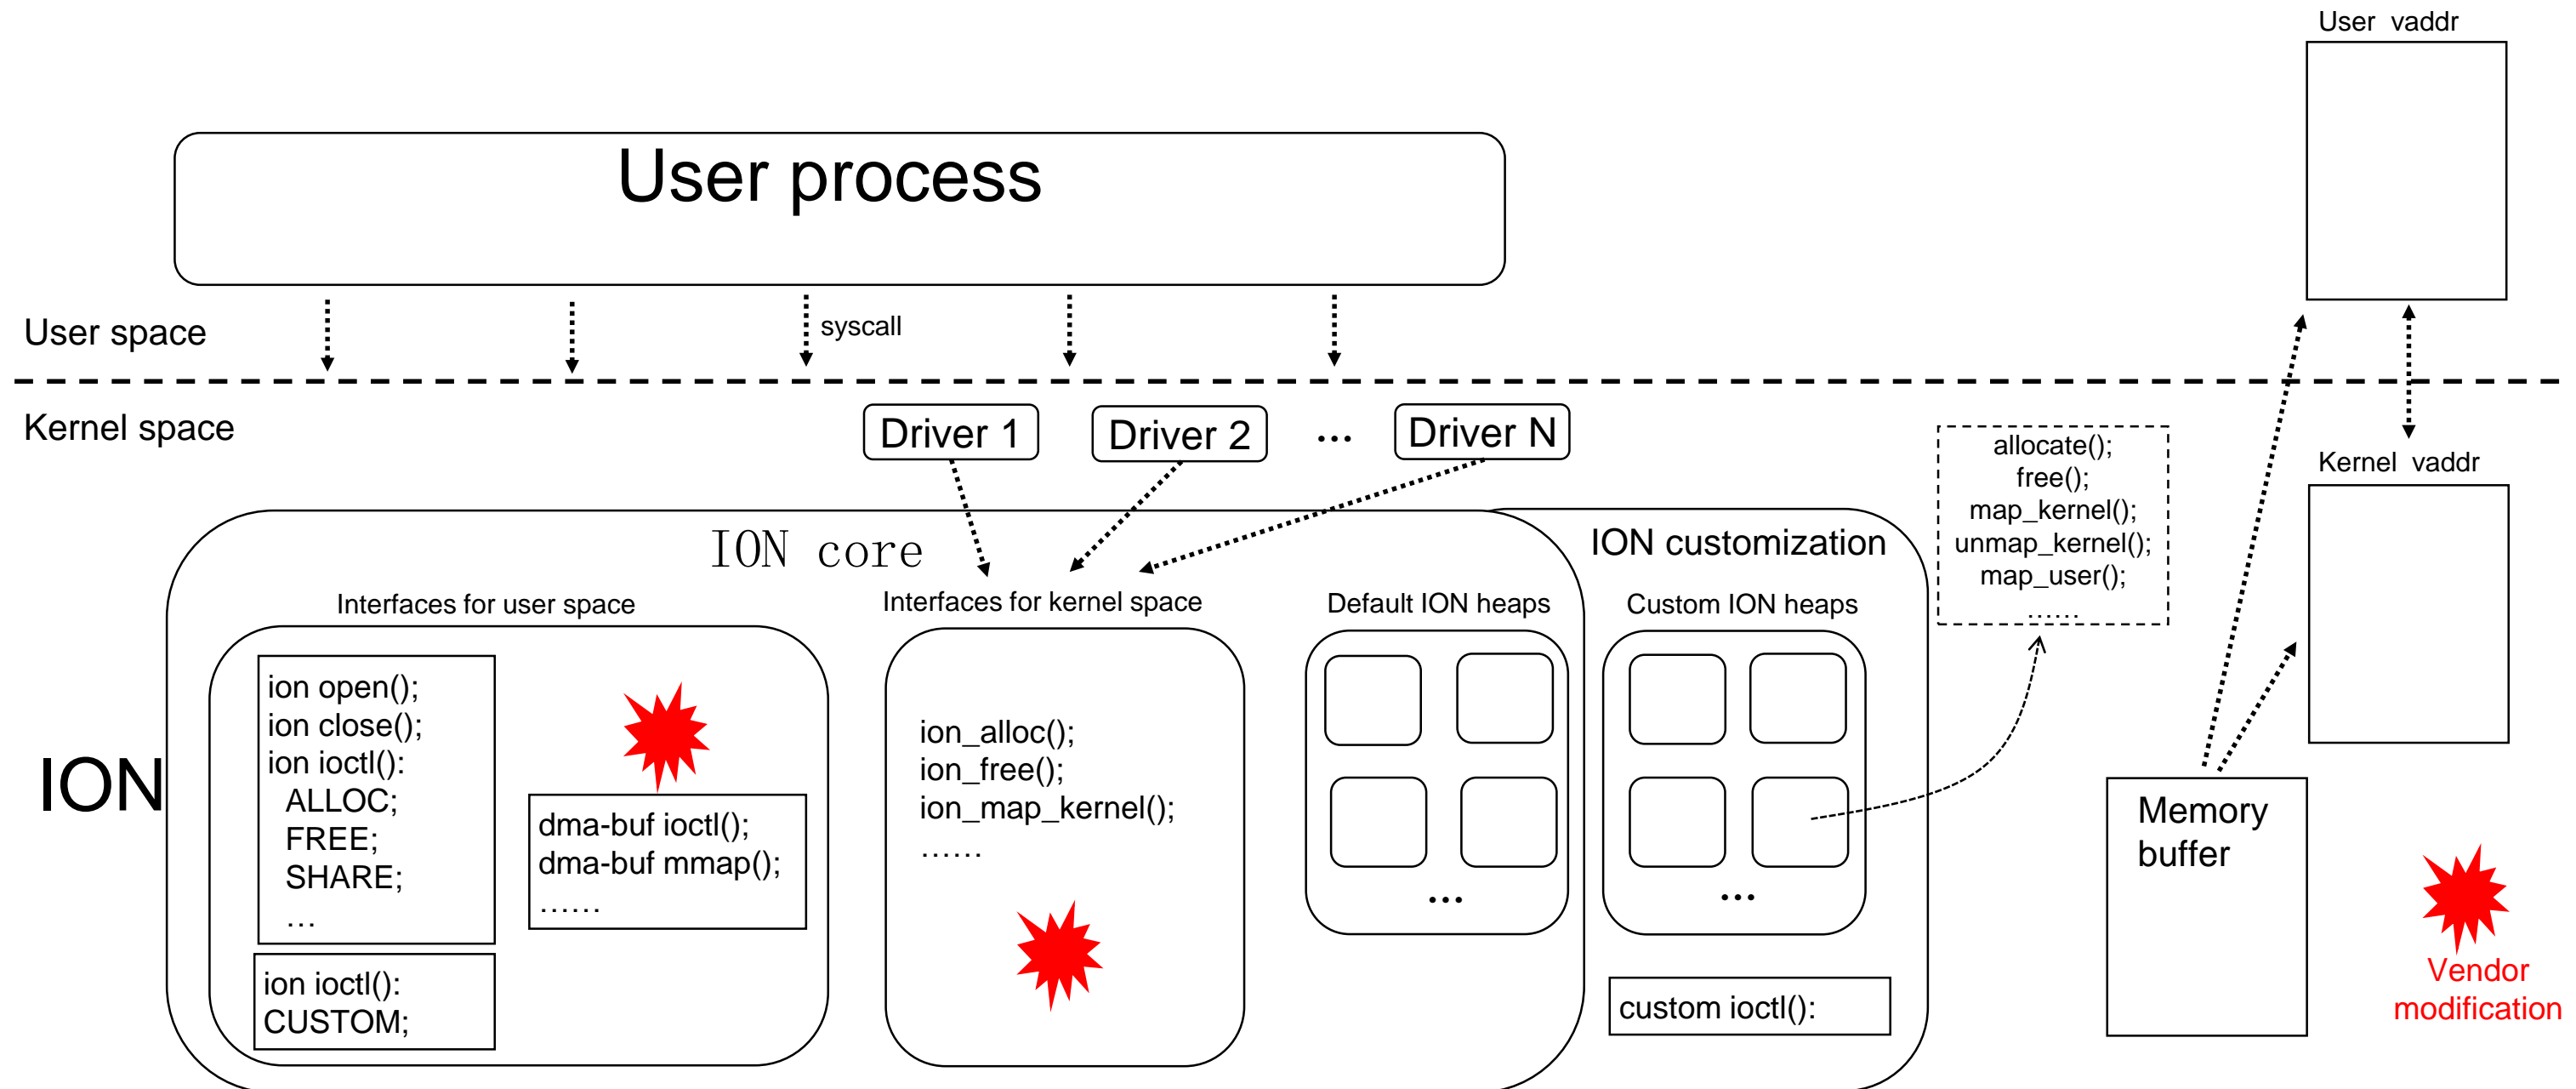
CVE-2022-20036,CVE-2021-0419,CVE-2021-0418,CVE-2021-0529,CVE-2021-0415,  
CVE-2021-0417,CVE-2021-0416,CVE-2022-20017,CVE-2022-20037,CVE-2021-0425



All the vulnerabilities have  
already been fixed by vendors!



# Vendors' modification to ION core





# Vendors' modification to ION core

## Vendor A

[Add new commands into the original ion\\_ioctl\(\):](#)

```
static long ion_ioctl(struct file *filp, unsigned int cmd,
unsigned long arg)
{
.....
+     case ION_IOC_CLEAN_CACHES:
+         return client->dev->custom_ioctl(client,
+
+         ION_IOC_CLEAN_CACHES, arg);
+     case ION_IOC_INV_CACHES:
+         return client->dev->custom_ioctl(client,
+
+         ION_IOC_INV_CACHES, arg);
+     case ION_IOC_CLEAN_INV_CACHES:
+         return client->dev->custom_ioctl(client,
+
+         ION_IOC_CLEAN_INV_CACHES, arg);
.....
```

## Vendor B

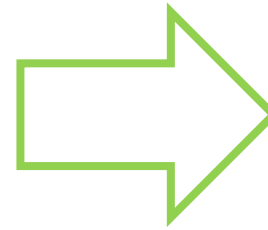
Add some debug code into the original ion\_alloc():

```
struct ion_handle *ion_alloc(struct ion_client *client, size_t len,
                             size_t align, unsigned int heap_id_mask,
                             unsigned int flags)
{
.....
+ #ifdef CONFIG_MTK_ION
+     ion_history_count_kick(true, len);
+ #endif

+     handle->dbg.user_ts = end;
+     do_div(handle->dbg.user_ts, 1000000);
+     memcpy(buffer->alloc_dbg, client->dbg_name,
ION_MM_DBG_NAME_LEN);
.....
```

## Vendors' modification to ION core


- Vulnerable modifications



New vulnerabilities  
can be introduced

- Code merge conflicts could happen

- Missing patch
- Wrong patch

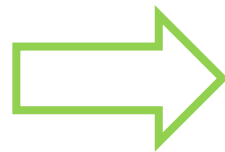





Known vulnerabilities are not  
fixed or new vulnerabilities  
get introduced



## Vendors' modification to ION core

- Analyze the modifications made by vendors
- Investigate known issues of ION and check if they are fixed downstream



- Vulnerabilities introduced by Modification  
CVE-2021-0422 
- Vulnerabilities due to the missing patch  
CVE-2021-39801   
CVE-2022-20109
- Vulnerabilities due to the wrong patch  
CVE-2022-20118   
CVE-2021-0610

# CVE-2021-0422

- Affected ION version:  
A vendor's devices with ION V1
- Root cause:  
Vulnerable modifications to ION core
- Impact:  
Illegal memory access

# CVE-2021-0422

```
struct ion_handle *ion_import_dma_buf_fd(struct ion_client
*client, int fd)
{
    struct dma_buf *dmabuf;
    struct ion_handle *handle;
    .....
    dmabuf = dma_buf_get(fd);
    .....
    handle = ion_import_dma_buf(client, dmabuf);
    dma_buf_put(dmabuf);
    .....
    handle->dbg.fd = fd;
    handle->dbg.user_ts = sched_clock();
    do_div(handle->dbg.user_ts, 1000000);
    return handle;
}
```

“handle” could be an error code

“handle” is not checked, invalid memory access will happen!



# CVE-2021-39801

- Affected ION version:  
ION V1
- Root cause:  
The patch of a known issue(wrong behavior of the ION API) is missing
- Impact:  
UAF

# CVE-2021-39801


## A known issue in ION\_IOC\_FREE:

```
case ION_IOC_FREE:
{
    struct ion_handle *handle;
    mutex_lock(&client->lock);
    handle = ion_handle_get_by_id_nolock(client,
                                         data.handle.handle);

    .....
    ion_free_nolock(client, handle);
    ion_handle_put_nolock(handle);
    mutex_unlock(&client->lock);
    break;
}
```

```
static void ion_handle_destroy(struct kref *kref)
{
    struct ion_handle *handle = container_of(kref, struct
    ion_handle, ref);
    .....
    kfree(handle);
}
```

```
handle_data.handle = ion_handle_id;
while(1) {
    ioctl(client_fd, ION_IOC_FREE,
    &handle_data);
}
```



**ion\_handle object in the kernel  
could be released at any time  
from user space!**


# CVE-2021-39801

UAF scenario:

```
case ION_IOC_ALLOC:
{
    struct ion_handle *handle;
    handle = ion_alloc(client, data.allocation.len,
                      data.allocation.align,
                      data.allocation.heap_id_mask,
                      data.allocation.flags);

    if (IS_ERR(handle))
        return PTR_ERR(handle);
    data.allocation.handle = handle->id;
    cleanup_handle = handle;
    break;
}
.....

if (copy_to_user((void __user *)arg, &data, _IOC_SIZE(cmd))) {
    if (cleanup_handle)
        ion_free(client, cleanup_handle);
    return -EFAULT;
}
```



Are there any  
other places that  
would trigger the  
UAF?

# CVE-2021-39801

The issue is assigned CVE-id: CVE-2017-0564 in 2017.

Patch for the issue:

[ANDROID: ion: Protect kref from userspace manipulation](#)

This separates the kref for ion handles into two components. Userspace requests through the ioctl will hold at most one reference to the internally used kref. All additional requests will increment a separate counter, and the original reference is only put once that counter hits 0. This protects the kernel from a poorly behaving userspace.

But some Android common kernel branches & Upstream kernel branches & some vendors' kernel branches missed the patch!

# CVE-2022-20118

- Affected ION version:  
A vendor's devices with ION V1
- Root cause:  
Wrong patch for a known mutex lock using issue
- Impact:  
UAF

# CVE-2022-20118

A known UAF issue in the ION\_IOC\_SHARE :

```
case ION_IOC_SHARE:
case ION_IOC_MAP:
{
    struct ion_handle *handle;
    handle = ion_handle_get_by_id(client, data.handle.handle);
    if (IS_ERR(handle))
        return PTR_ERR(handle);
    data.fd.fd = ion_share_dma_buf_fd(client, handle);
    ion_handle_put(handle);
    if (data.fd.fd < 0)
        ret = data.fd.fd;
    break;
}
```

```
case ION_IOC_FREE:
{
    struct ion_handle *handle;
    mutex_lock(&client->lock);
    handle = ion_handle_get_by_id_nolock(client,
data.handle.handle);
    if (IS_ERR(handle)) {
        mutex_unlock(&client->lock);
        return PTR_ERR(handle);
    }
    ion_free_nolock(client, handle);
    ion_handle_put_nolock(handle);
    mutex_unlock(&client->lock);
    break;
}
```

Wrong behavior  
version of  
ION\_IOC\_FREE

# CVE-2022-20118

A known UAF issue in the ION\_IOC\_SHARE:

## Thread A

(Process ION\_IOC\_SHARE / ION\_IOC\_MAP)

```
handle = ion_handle_get_by_id(client,  
data.handle.handle);
```

```
data.fd.fd = ion_share_dma_buf_fd(client, handle);
```

UAF

## Thread B

(user space)

```
handle_data.handle = ion_handle_id;  
  
ioctl(client_fd, ION_IOC_FREE, &handle_data);  
ioctl(client_fd, ION_IOC_FREE, &handle_data);  
(the ion_handle object in the kernel will be released  
after the two ION_IOC_FREE)
```

# CVE-2022-20118

The [patch](#) to fix the issue in 2018:

staging: android: ion: fix ION\_IOC\_{MAP,SHARE} use-after-free:

```
case ION_IOC_SHARE:
case ION_IOC_MAP:
{
    struct ion_handle *handle;
    mutex_lock(&client->lock);
    handle = ion_handle_get_by_id_nolock(client, data.handle.handle);
    if (IS_ERR(handle)) {
        mutex_unlock(&client->lock);
        return PTR_ERR(handle);
    }
    data.fd.fd = ion_share_dma_buf_fd_nolock(client, handle);
    ion_handle_put_nolock(handle);
    mutex_unlock(&client->lock);
    if (data.fd.fd < 0)
        ret = data.fd.fd;
    break;
}
```

```
case ION_IOC_FREE:
{
    struct ion_handle *handle;
    mutex_lock(&client->lock);
    handle = ion_handle_get_by_id_nolock(client,
data.handle.handle);
    if (IS_ERR(handle)) {
        mutex_unlock(&client->lock);
        return PTR_ERR(handle);
    }
    ion_free_nolock(client, handle);
    ion_handle_put_nolock(handle);
    mutex_unlock(&client->lock);
    break;
}
```



# CVE-2022-20118

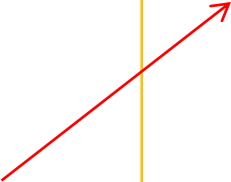
A variant vulnerability similar to the known issue:

In a vendor's ION driver:

```
case ION_IOC_SHARE:
case ION_IOC_MAP:
{
    struct ion_handle *handle;

    handle = ion_handle_get_by_id_nolock(client,
data.handle.handle);
    if (IS_ERR(handle))
        return PTR_ERR(handle);
    data.fd.fd = ion_share_dma_buf_fd(client, handle);
    ion_handle_put(handle);
    if (data.fd.fd < 0)
        ret = data.fd.fd;

    break;
}
```



```
struct ion_handle *ion_handle_get_by_id_nolock(struct ion_client
*client,
                                             int id)
{
    struct ion_handle *handle;
    handle = idr_find(&client->idr, id);
    if (handle)
        ion_handle_get(handle);
    return handle ? handle : ERR_PTR(-EINVAL);
}
```

```
case ION_IOC_FREE:
(correct behavior version)
```

# CVE-2022-20118

UAF still happens in function `ion_handle_get_by_id_nolock`:

## Thread A

(Enter function `ion_handle_get_by_id_nolock`)

```
handle = idr_find(&client->idr, id);
```

```
if (handle)
```

```
ion_handle_get(handle)
```

**UAF**

## Thread B

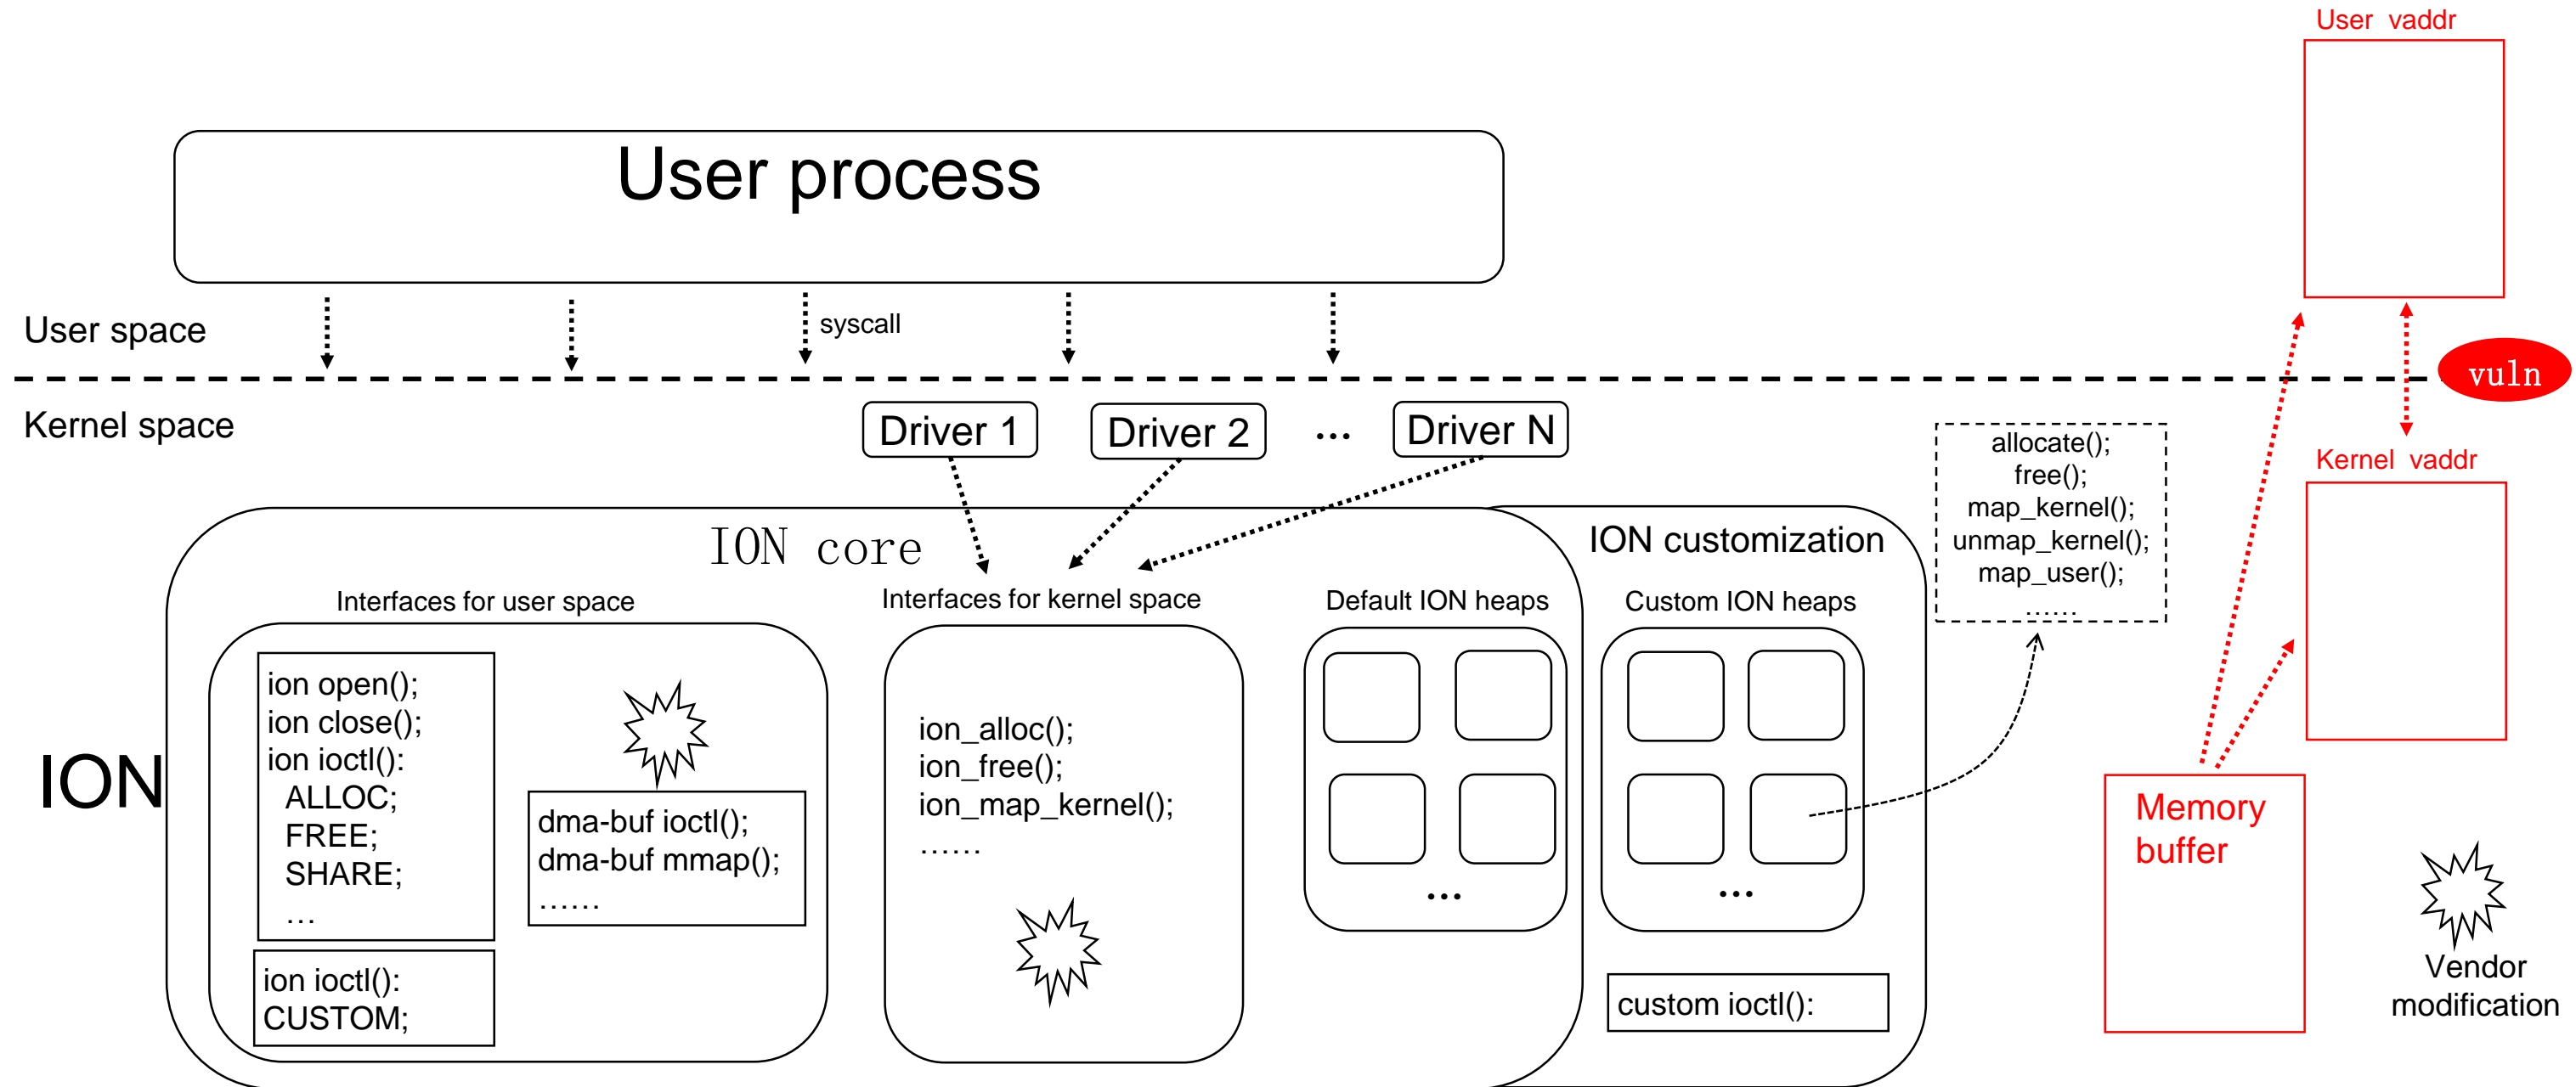
(User space)

```
handle_data.handle = ion_handle_id;
```

```
ioctl(client_fd, ION_IOC_FREE, &handle_data); (the  
ion_handle object in the kernel will be released after the  
ION_IOC_FREE)
```



# Buffer sharing



# Buffer sharing feature of ION can introduce vulnerabilities:

## ➤ 《[Android ION Hazard](#)》

Two kinds of vulnerabilities introduced by ION buffering sharing

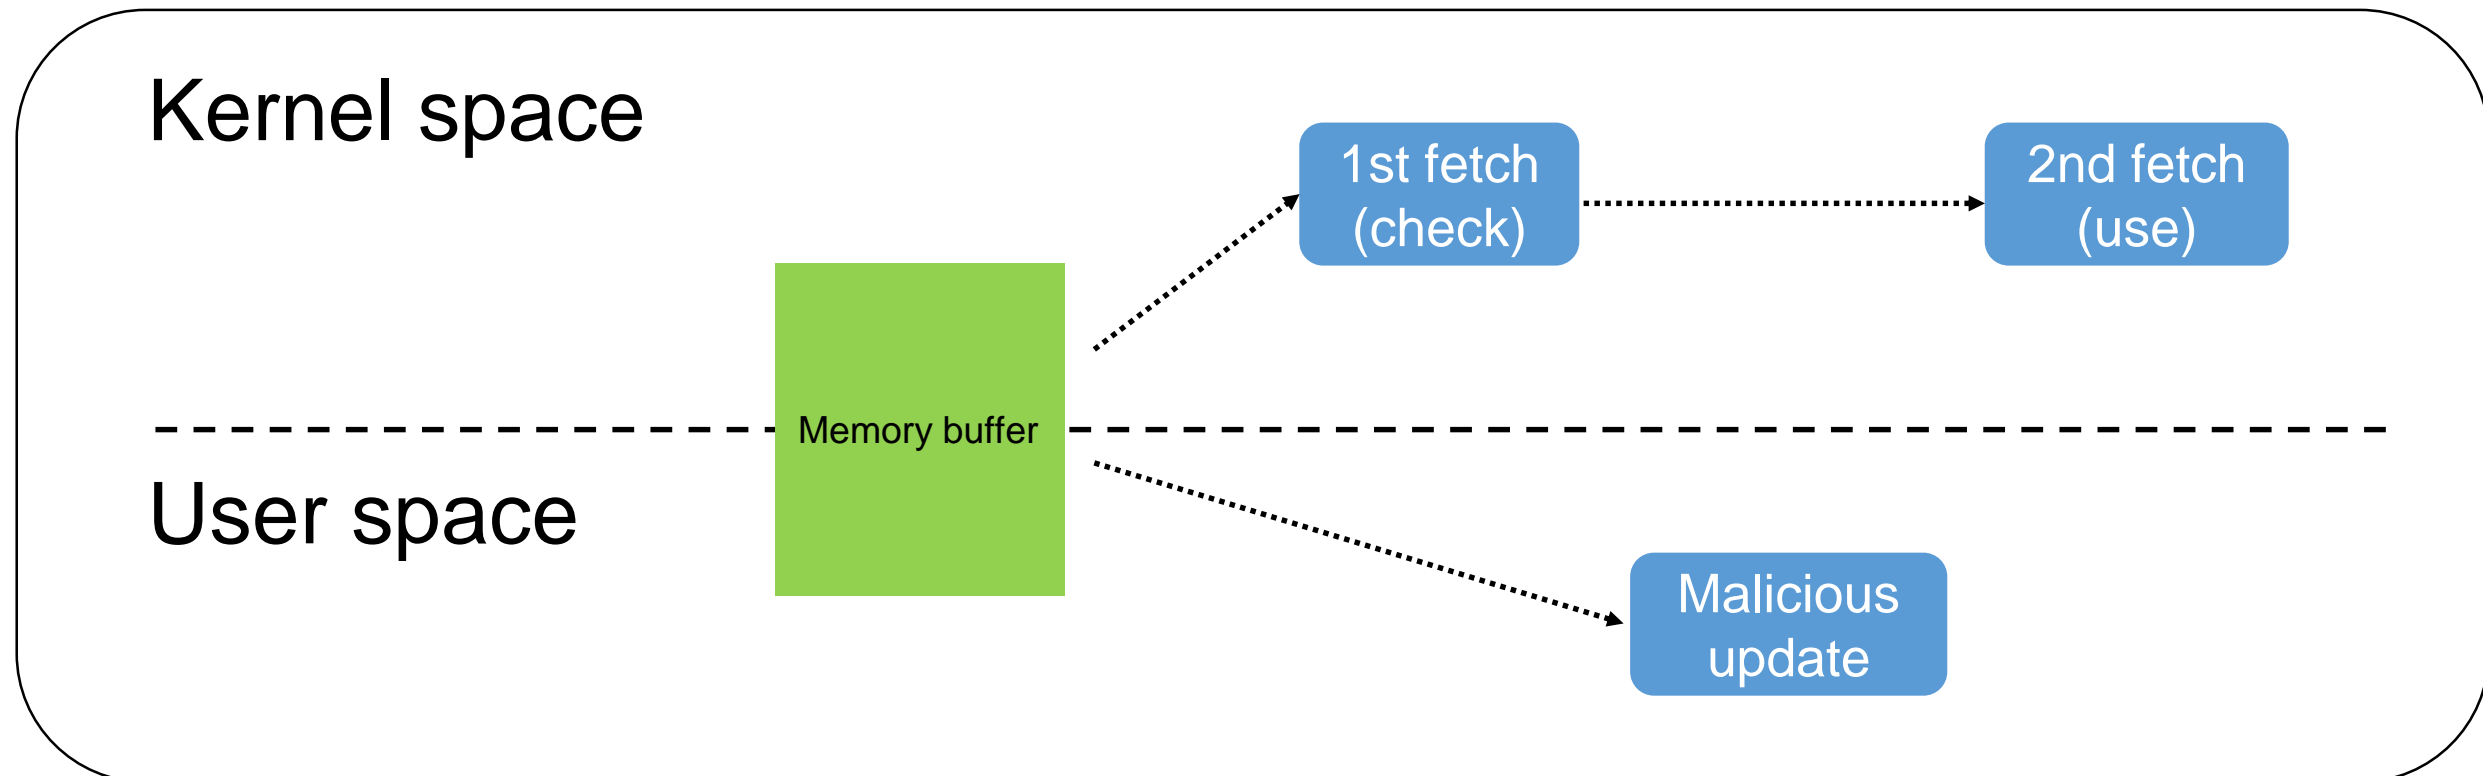
- System crash due to hardware protection

- Sensitive information leakage

# Buffer sharing feature of ION can introduce vulnerabilities:

## ➤ 《[An iOS hacker tries Android](#)》

- Double fetch vulnerabilities introduced by Buffer Sharing



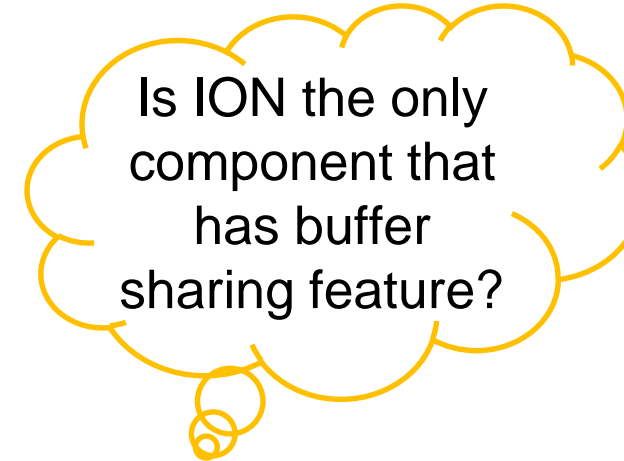
## Double fetch vulnerabilities introduced by Buffer Sharing

We found other double fetch vulnerabilities in a vendor's apu driver:

**CVE-2021-0897**

**CVE-2021-0895**

**CVE-2021-0903**



# Exploitation

# Get root!

```

下午4:33
窗口 1
+ X ...

./ $
./ $ getprop ro.build.fingerprint
Redmi/atom/atom:10/QP1A.190711.020/V12.0.5.0.QJHCNXM:user/release-keys
./ $ getenforce
getenforce: Couldn't get enforcing status: Permission denied
1|./ $ id
uid=10223(u0_a223) gid=10223(u0_a223) groups=10223(u0_a223),3003(inet),9997(everybody),20223(u0_a223_cache),50223(all_a223) context=u:r:untrusted_app_25:s0:c512,c768
./ $ ~/explode

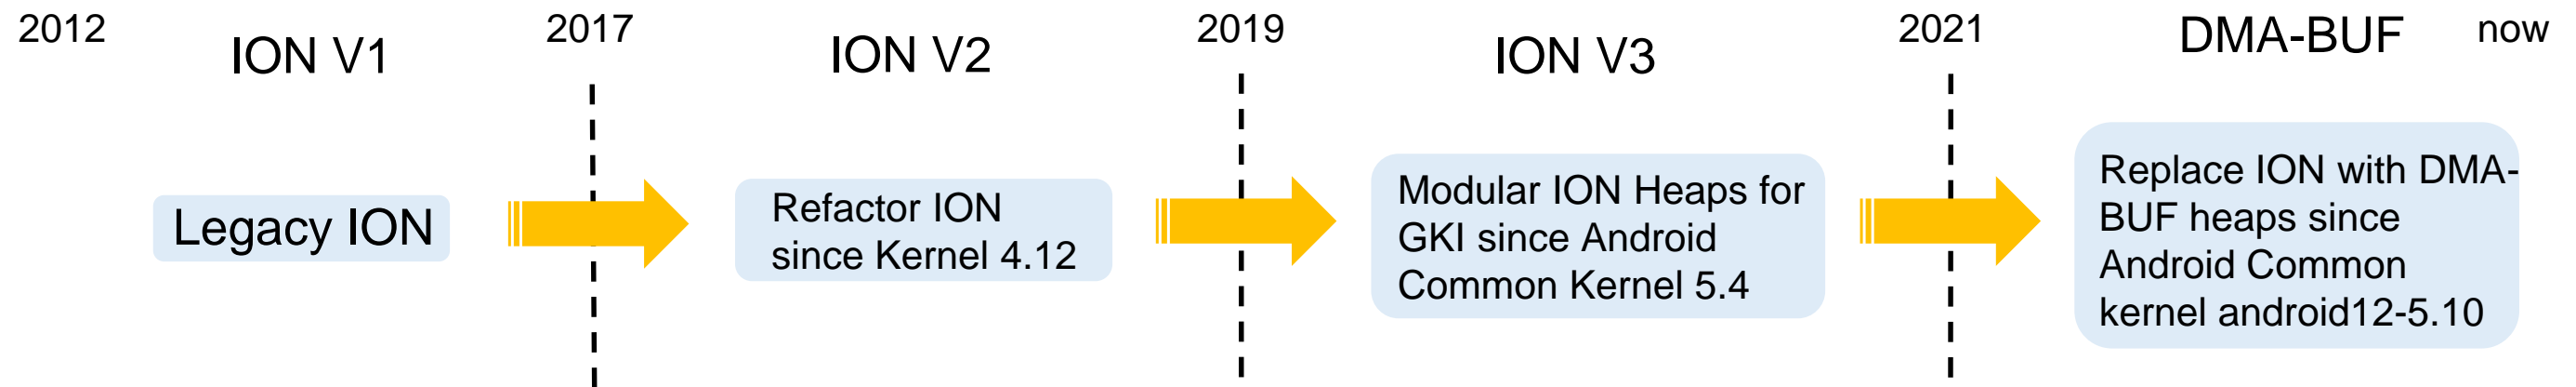
#####
#                                     #
#               *       *       *       #
#             *   *   *   *       #
#           * * * * *ExplosION * * * * * #
#         * * * * *       *       #
#       *   *   *   *       #
#     *       *       *       #
#                                     #
#       Exploited by Le Wu of Baidu Security       #
#                                     #
#####

[+] Trigger ...
[+] Bypass KASLR
kaslr_slide:0x1732800000
[+] SELinux disabled
[+] Credential altered
[+] Pwn!
./ # id
uid=0(root) gid=0(root) groups=0(root),3003(inet),9997(everybody),20223(u0_a223_cache),50223(all_a223) context=u:r:untrusted_app_25:s0:c512,c768
./ # getenforce
Permissive
./ #

```

# Reflections on ExplosION

## ➤ ION Evolution——Refactors from Upstream

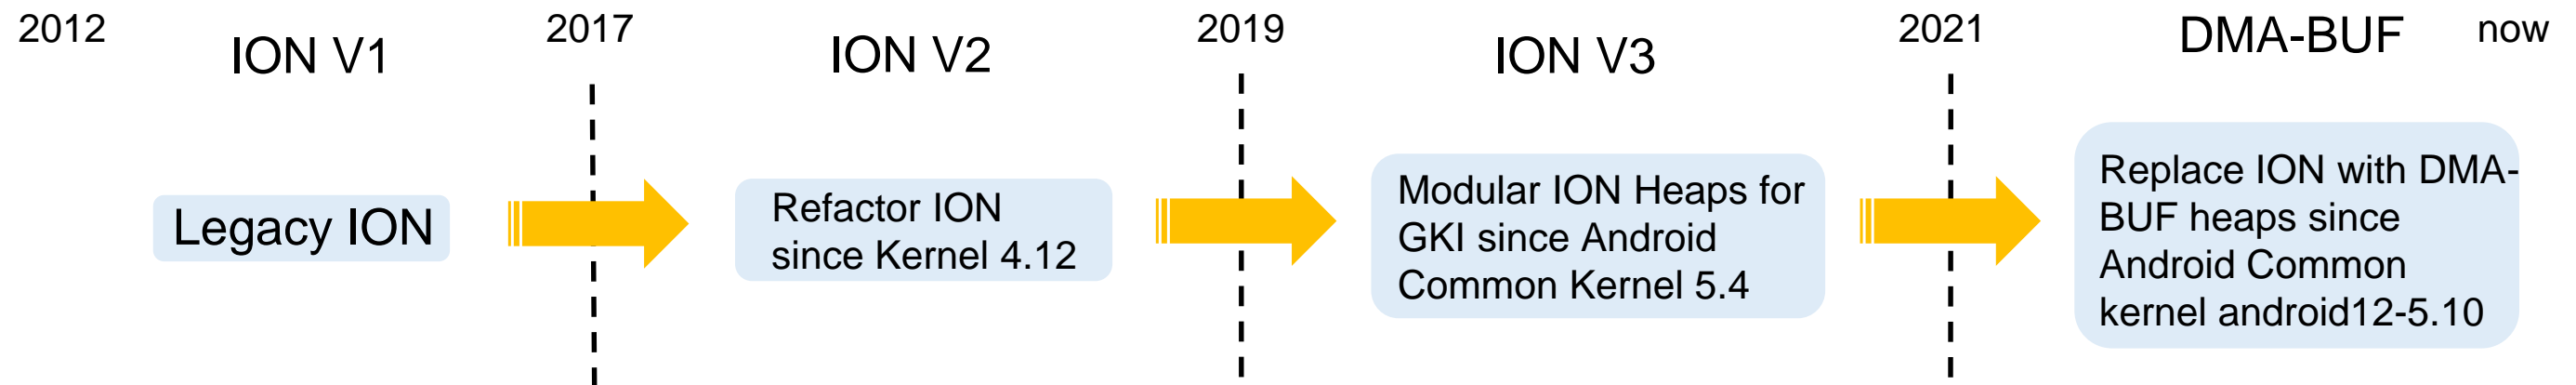


- Have been used since around 2012
- Still widely used in devices of which kernel < 4.12
- Bug fixes & small refactors happen from time to time



# Reflections on ExplosION

## ➤ ION Evolution——Refactors from Upstream

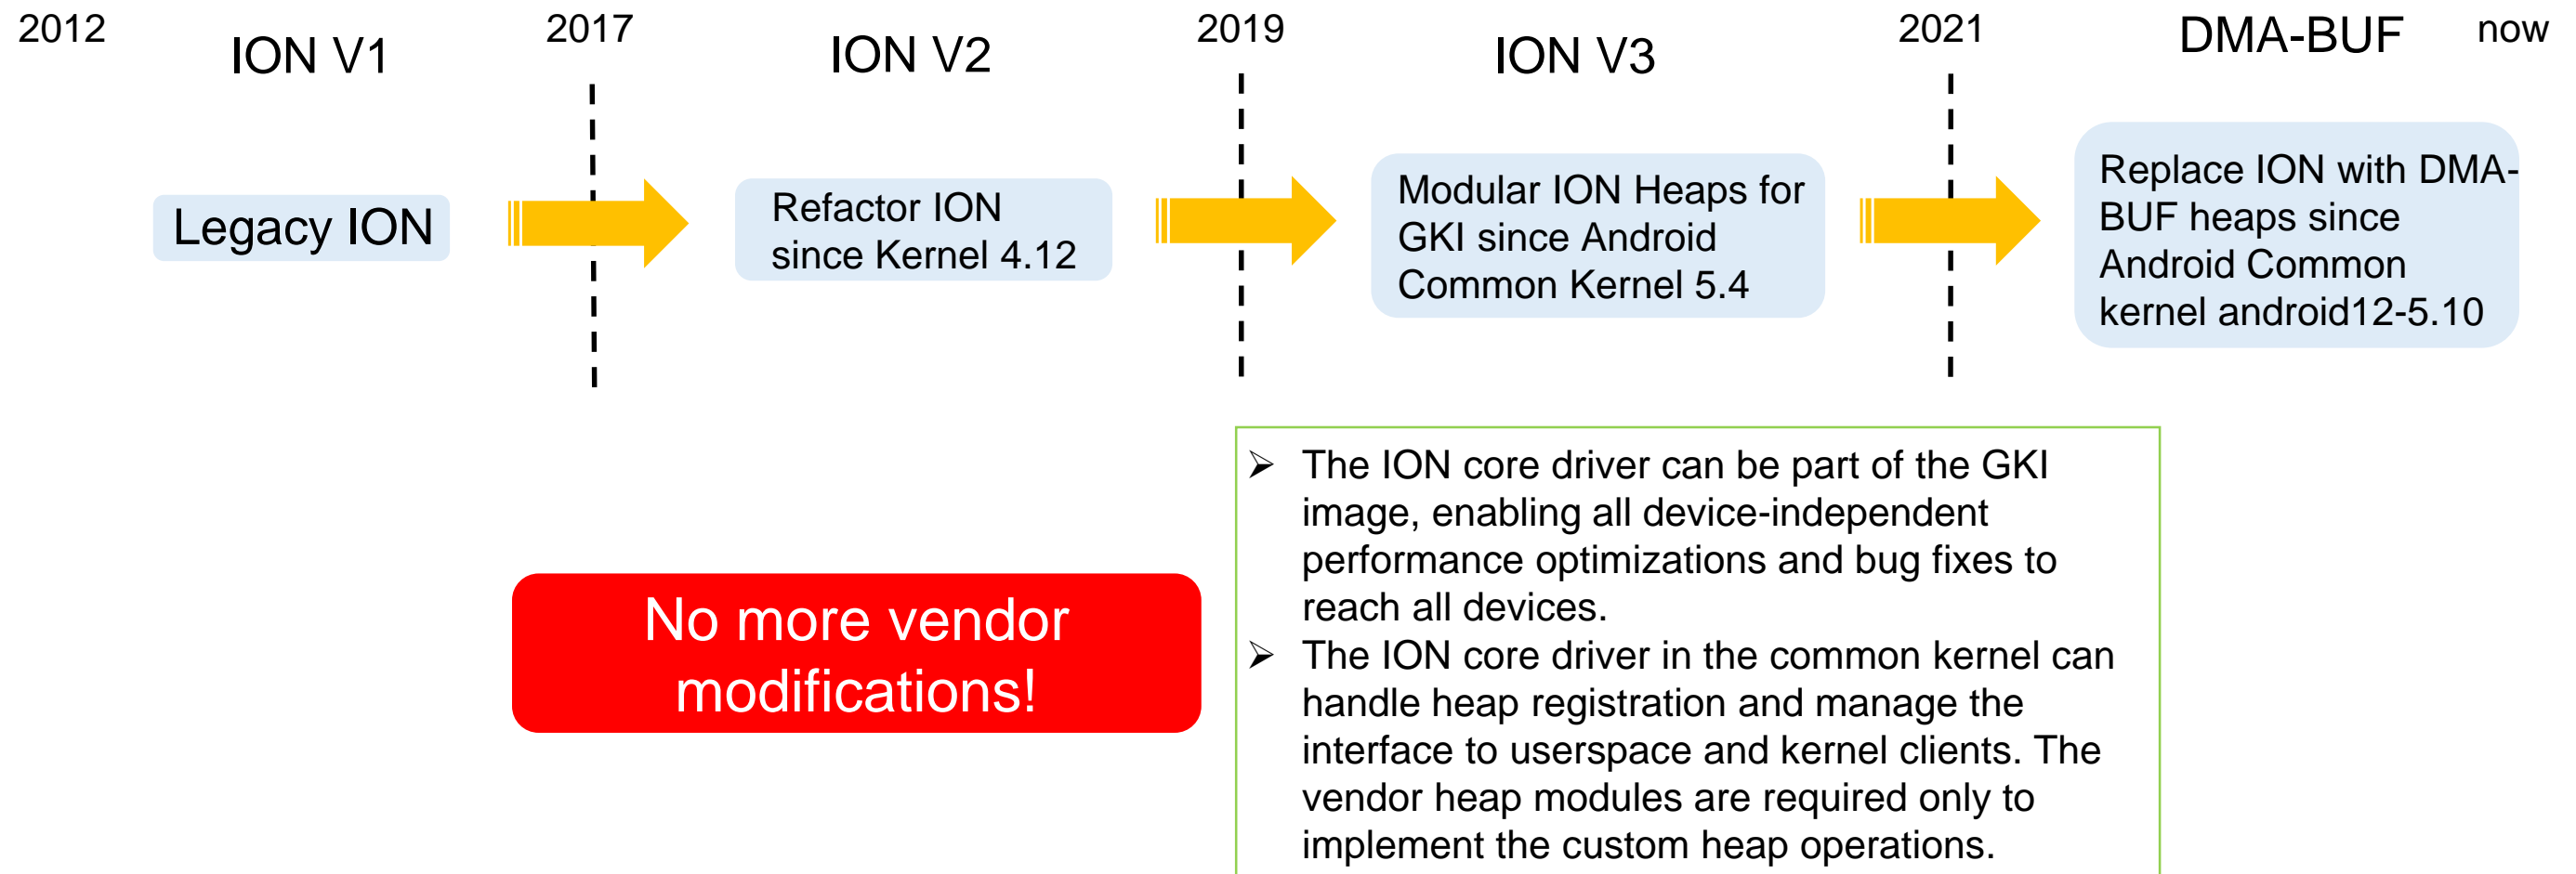


- Removal of Ion clients and handles: `IOC_ION_ALLOC` ioctl directly outputs dma-buf fds.
- Addition of cache-coherency ioctls: Kernel 4.12 replaced `ION_IOC_SYNC` with the `DMA_BUF_IOCTL_SYNC` ioctl

ION core become more brief, which means more secure!

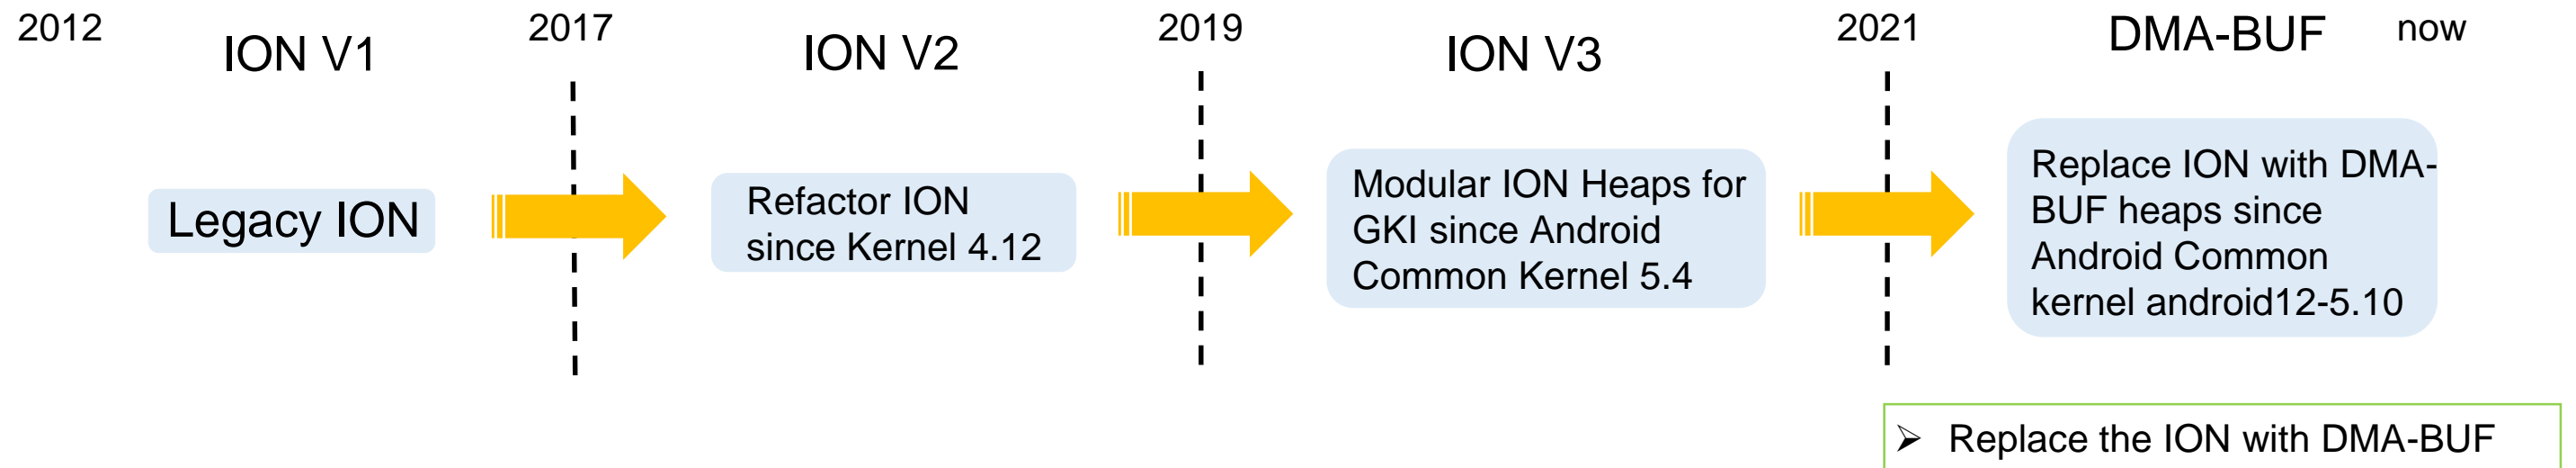
# Reflections on ExplosION

## ➤ ION Evolution——Refactors from Upstream



# Reflections on ExplosION

## ➤ ION Evolution——Refactors from Upstream



Advantages of DMA-BUF heaps:

- ✓ Security
- ✓ ABI stability
- ✓ Standardization

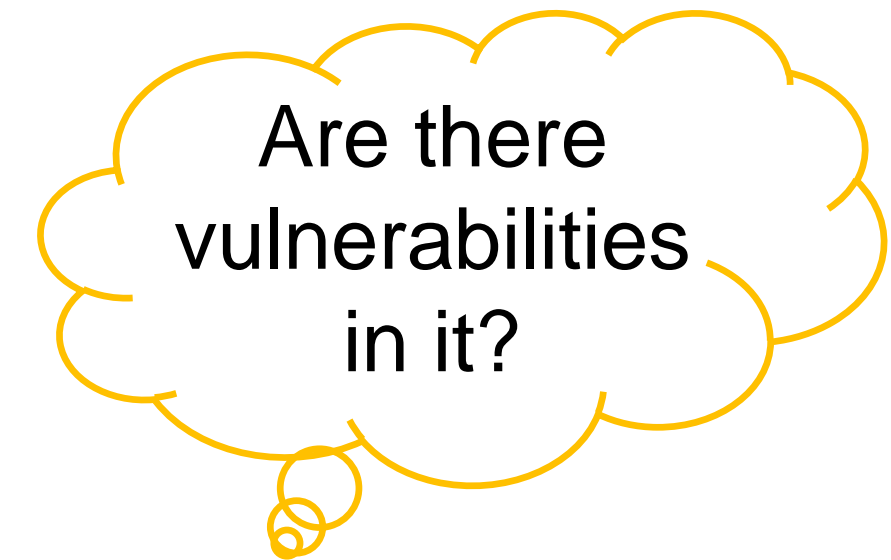
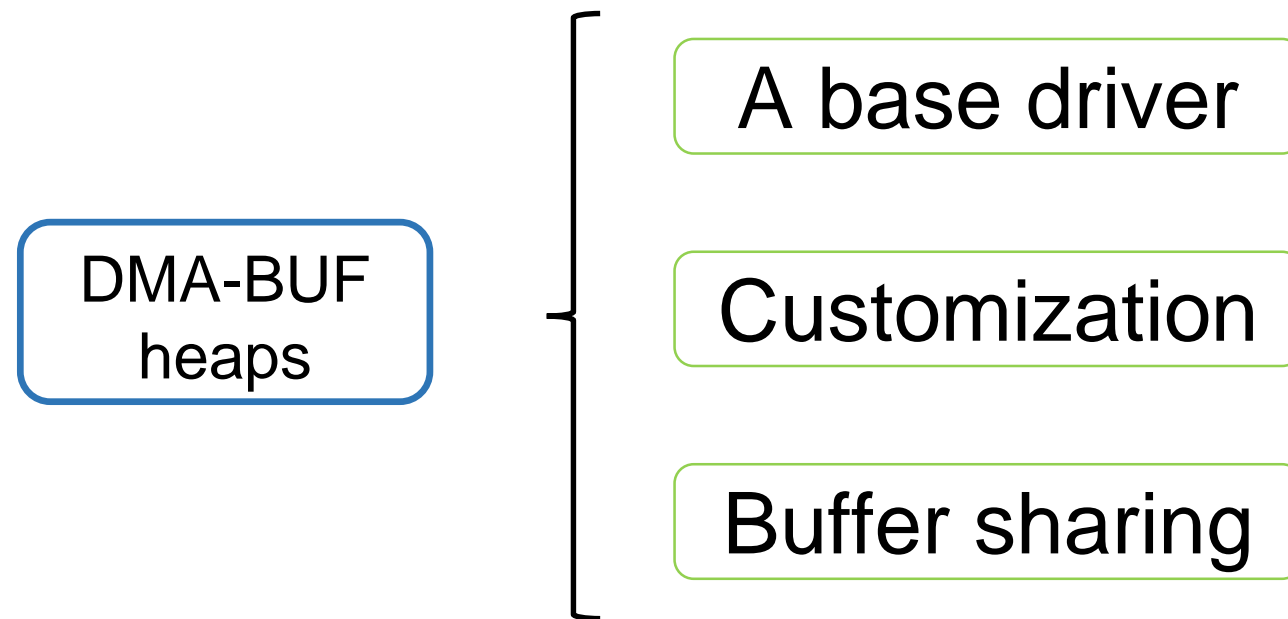
# Reflections on ExplosION

## ➤ Suggestions for vendors

- Apply the refactors right now!
- Patch quickly and carefully!
- Understand the ION APIs before using them!
- Do customization carefully!

All the affected vendors had worked diligently with us to remediate the ExplosION and had already made the patches available.

# Future work



# Acknowledge

- Thanks to Chengfu Bao, 某因幡, Shufan Yang, Lin Wu

**Thank you!**

# Supplement

## **Details of other ION vulnerabilities**



# CVE-2021-39714

- Affected ION version:  
ION V1,V2
- Root cause:  
Reference count overflow in ION core
- Impact:  
Potential UAF would be introduced into vendor drivers  
by ION

# CVE-2021-39714

```
void *ion_map_kernel(struct ion_client *client, struct
ion_handle *handle)
{
    struct ion_buffer *buffer;
    void *vaddr;
    mutex_lock(&client->lock);
    .....
    buffer = handle->buffer;
    if (!handle->buffer->heap->ops->map_kernel) {
        .....
    }
    mutex_lock(&buffer->lock);
    vaddr = ion_handle_kmap_get(handle);
    mutex_unlock(&buffer->lock);
    mutex_unlock(&client->lock);
    return vaddr;
}
```

```
static void *ion_buffer_kmap_get(struct ion_buffer *buffer)
{
    void *vaddr;
    if (buffer->kmap_cnt) {
        buffer->kmap_cnt++;
        return buffer->vaddr;
    }
    vaddr = buffer->heap->ops->map_kernel(buffer-
>heap, buffer);
    .....
    buffer->vaddr = vaddr;
    buffer->kmap_cnt++;
    return vaddr;
}
```

Integer overflow can happen !!!

# CVE-2021-39714

```
void ion_unmap_kernel(struct ion_client *client, struct
ion_handle *handle)
{
    struct ion_buffer *buffer;
    mutex_lock(&client->lock);
    buffer = handle->buffer;
    mutex_lock(&buffer->lock);
    ion_handle_kmap_put(handle);
    mutex_unlock(&buffer->lock);
    mutex_unlock(&client->lock);
}
```

```
static void ion_buffer_kmap_put(struct ion_buffer *buffer)
{
    buffer->kmap_cnt--;
    if (!buffer->kmap_cnt) {
        buffer->heap->ops->unmap_kernel(buffer->heap,
buffer);
        buffer->vaddr = NULL;
    }
}
```

# CVE-2021-39714

## Thread A

call `ion_map_kernel()` constantly to let `ion_buffer->kmap_cnt` become `0xffffffff`

`vaddr2 = ion_map_kernel();`  
`ion_buffer->kmap_cnt` become 1;

Access the `vaddr2`;

**UAF**

## Thread B

`vaddr = ion_map_kernel();`  
`ion_buffer->kmap_cnt` become 0; Integer overflow

`ion_unmap_kernel();`  
`vaddr2` will be unmapped

# CVE-2022-20109

- Affected ION version:  
ION V1

- Root cause:

The patch of a known issue is missing, resulting in a reference count issue of ion\_handle

- Impact:  
UAF

# CVE-2022-20109

The latest legacy ION\_IOC\_ALLOC:

```
case ION_IOC_ALLOC:
{
    struct ion_handle *handle;
    handle = __ion_alloc(client, data.allocation.len,
                        data.allocation.align,
                        data.allocation.heap_id_mask,
                        data.allocation.flags, true);

    .....
    cleanup_handle = handle;
    pass_to_user(handle);
    break;
}
.....
if (copy_to_user((void __user *)arg, &data, _IOC_SIZE(cmd))) {
    if (cleanup_handle) {
        mutex_lock(&client->lock);
        user_ion_free_nolock(client, cleanup_handle);
        ion_handle_put_nolock(cleanup_handle);
        mutex_unlock(&client->lock);
    }
    .....
}
```

Reference count of  
ion\_handle becomes 2

Reference count of  
ion\_handle becomes 1

Reference count of ion\_handle  
becomes 0, ion\_handle will be  
released!

# CVE-2022-20109

ION\_IOC\_ALLOC of a vendor's devices is like this:

```
case ION_IOC_ALLOC:
```

```
{
```

```
    struct ion_handle *handle;
```

```
    handle = ion_alloc(client, data.allocation.len,  
                      data.allocation.align,  
                      data.allocation.heap_id_mask,  
                      data.allocation.flags);
```

Reference count of  
ion\_handle becomes 1

```
    .....
```

```
    pass_to_user(handle);  
    data.allocation.handle = handle->id;  
    cleanup_handle = handle;
```

```
    .....
```

```
    if (copy_to_user((void __user *)arg, &data, _IOC_SIZE(cmd))) {  
        if (cleanup_handle) {
```

```
            mutex_lock(&client->lock);  
            user_ion_free_nolock(client, cleanup_handle);  
            ion_handle_put_nolock(cleanup_handle);
```

Reference count of  
ion\_handle becomes  
0, ion\_handle will be  
released!

```
    .....
```

**UAF**

# CVE-2022-20109

The root cause of the UAF: Missed a [patch](#) which is released in **2016**

ion: Fix use after free during ION\_IOC\_ALLOC

If a user happens to call ION\_IOC\_FREE during an ION\_IOC\_ALLOC on the just allocated id, and the copy\_to\_user fails, the cleanup code will attempt to free an already freed handle.

This adds a wrapper for ion\_alloc that adds an ion\_handle\_get to avoid this.



# CVE-2021-0610

- Affected ION version:  
ION V1

- Root cause:  
The patch of a known issue(ion\_handle kref overflow) is missing

- Impact:

# CVE-2021-0610

A known issue 3 years ago:

```
case ION_IOC_IMPORT:
{
    struct ion_handle *handle;
    handle = ion_import_dma_buf_fd(client, data.fd.fd);
    if (IS_ERR(handle)) {
        ret = PTR_ERR(handle);
    } else {
        data.handle.handle = handle->id;
        handle = pass_to_user(handle);
        if (IS_ERR(handle)) {
            ret = PTR_ERR(handle);
            data.handle.handle = 0;
        }
    }
    break;
}
```

```
struct ion_handle *ion_import_dma_buf(struct ion_client *client,
                                      struct dma_buf *dmabuf)
{
    struct ion_buffer *buffer;
    struct ion_handle *handle;
    int ret;
    .....
    buffer = dmabuf->priv;
    mutex_lock(&client->lock);
    /* if a handle exists for this buffer just take a reference to it */
    handle = ion_handle_lookup(client, buffer);
    if (!IS_ERR(handle)) {
        ion_handle_get(handle);
        mutex_unlock(&client->lock);
        goto end;
    }
    .....
}
```

kref overflow !

# CVE-2021-0610

## The patch to fix it:

staging: android: ion: check for kref overflow

This patch is against 4.9. It does not apply to master due to a large rework of ion in 4.12 which removed the affected functions altogether.

4c23cbff073f3b9b ("staging: android: ion: Remove import interface")

Userspace can cause the kref to handles to increment arbitrarily high. Ensure it does not overflow.

```
+/* Must hold the client lock */
+static struct ion_handle *ion_handle_get_check_overflow(
+    struct ion_handle *handle)
+{
+    if (atomic_read(&handle->ref.refcount) + 1 == 0)
+        return ERR_PTR(-E_OVERFLOW);
+    ion_handle_get(handle);
+    return handle;
+}
+
+.....
static bool ion_handle_validate(struct ion_client *client,
@@ -1110,7 +1121,7 @@
    /* if a handle exists for this buffer just take a reference to it */
    handle = ion_handle_lookup(client, buffer);
    if (!IS_ERR(handle)) {
-        ion_handle_get(handle);
+        handle = ion_handle_get_check_overflow(handle);
        mutex_unlock(&client->lock);
        goto end;
    }
```

# CVE-2021-0610

The issue still exists in a vendor's devices because of the wrong patch:

```
struct ion_handle *ion_import_dma_buf(struct ion_client *client,
                                     struct dma_buf *dmabuf)
{
    struct ion_buffer *buffer;
    struct ion_handle *handle;
    .....
    buffer = dmabuf->priv;

    mutex_lock(&client->lock);
    /* if a handle exists for this buffer just take a reference to it */
    handle = ion_handle_lookup(client, buffer);
    if (!IS_ERR(handle)) {
        ion_handle_get_check_overflow(handle);
        mutex_unlock(&client->lock);
        goto end;
    }
    .....
}
```

Kref overflow can still happen!!!

```
static struct ion_handle *ion_handle_get_check_overflow(
    struct ion_handle *handle)
{
    if (atomic_read(&handle->ref.refcount.refs) + 1 ==
        0)
        return ERR_PTR(-E_OVERFLOW);
    ion_handle_get(handle);
    return handle;
}
```